

Neural Representation of Color in Color Vision Deficiency: An fMRI Study

Your Name¹

¹Your University

Abstract

Placeholder abstract. Replace this with your actual abstract text. This study investigated neural color representations in individuals with color vision deficiency (CVD) using functional MRI and computational modeling approaches.

1 Introduction

Color vision deficiency (CVD) affects approximately 8% of males and 0.5% of females of Northern European descent. Despite its prevalence, the neural mechanisms underlying altered color perception in CVD remain incompletely understood. While genetic and psychophysical studies have characterized the photoreceptor-level basis of CVD, how these peripheral differences cascade into cortical color representations remains an open question.

Previous neuroimaging studies have demonstrated that color information is represented throughout the visual hierarchy, from early visual cortex (V1, V2) to higher-order color-selective regions (V4, VO1). Brouwer (2009) pioneered the use of forward encoding models to decode and reconstruct novel colors from V4 activity patterns, demonstrating that neural populations encode a continuous representation of hue space. This approach was extended by Brouwer (2013), who showed categorical clustering of color representations in ventral temporal cortex. More recently, multivariate pattern analysis has revealed that color information can be decoded from distributed patterns across the visual hierarchy (Haxby et al., 2011). However, most studies have focused on individuals with typical color vision, leaving open questions about how CVD alters these neural representations.

While behavioral studies have characterized perceptual differences in CVD (Brettel, Viénot & Mollon, 1997), the neural basis of these alterations has not been systematically investigated using multivariate pattern analysis and computational modeling approaches. Critically, it remains unknown whether CVD affects only the discriminability of individual colors or also distorts the geometric structure of neural color space. Specifically,

three questions remain unresolved:

1. Do neural color representations in CVD differ from healthy controls at the population level, and can these differences be characterized in a common representational space?
2. Is color discrimination (distinguishing trained colors) preserved while color interpolation (generalizing to novel colors) is impaired in CVD?
3. Can forward encoding models predict and correct for CVD-related distortions in neural color space?

The present study addresses these questions by combining functional magnetic resonance imaging (fMRI) with shared response modeling (Chen et al., 2015)—a dimension-reduction technique that enables comparison of neural representations across individuals—and forward encoding models (Brouwer & Heeger, 2009). We tested the hypothesis that CVD participants would show systematic deviations in neural color representations compared to healthy controls, particularly in color-selective regions of visual cortex. Crucially, we distinguished between discrimination (distinguishing trained colors) and interpolation (generalizing to novel colors) using leave-one-color-out cross-validation.

Understanding the neural basis of CVD has both theoretical and practical implications. Theoretically, it provides insights into how genetic variation in photoreceptor spectral sensitivity shapes cortical color representations, addressing long-standing questions about the relationship between peripheral and central color coding. Practically, characterizing neural distortions in CVD could inform the development of personalized color correction filters that restore not only color

discrimination but also the geometric structure of perceptual color space.

2 Data collection

Participants

Twelve volunteers were recruited under Institutional Review Board approval (#2510/002-023), providing written informed consent prior to participation. Participants reported no history of neurological disorders. Color vision was assessed using the 14-plate Ishihara test (Ishihara, 1917). Participants accurately identifying more than 12 plates were classified as Healthy Controls (HC; $n = 7$, 3 females, age 22.7 ± 2.5 years), while those scoring below this threshold were classified as having Color Vision Deficiency (CVD; $n = 2$, including one participant with deuteranomaly and one with protanomaly, both male). Three participants were excluded for failing to complete the experiment. Given the small CVD sample, all CVD results are interpreted as individual case demonstrations using Crawford and Howell (1998) single-case statistics, not as population-level effects.

fMRI experiment

We conducted an fMRI experiment to collect participants' voxel responses while perceiving colors. Stimuli consisted of eight colors equally spaced from 0° to 315° in Commission Internationale de L'Eclairage (CIE) $L^*a^*b^*$ space (Commission Internationale de L'Eclairage, 1986). A gray filler was additionally included as a nonchromatic condition. Their lightness was fixed to $L^* = 75$ with an identical distance of 40 chroma from the gray point with $a^* = 0$, and $b^* = 0$ (Fig 1A).

To maintain participants' attention and fixation, we conducted a modified 'Rapid Serial Visual Presentation' (RSVP) detection task (Brouwer & Heeger, 2009) throughout each run. During the experiment, a circular patch was presented for 1.5 s. The patch contained one of the color stimuli or a gray point in a randomized order. The order of colors was optimized using the Neurodesign package (Durnez, Moerkerke & Nichols, 2018) to maximize design efficiency. 0.35s after stimulus onset, a set of alternating alphabets was presented for 0.4s each above the stimulus. Participants were instructed to press a button when an alphabet 'K' appeared consecutively. Such a target set emerged with a random probability of 0.33. Interstimulus intervals (ISIs) were randomly selected among 3s, 4.5s, and 6s. Each run included eight repetitions of the eight colors and a gray point. Each run lasted approximately 7min and each session consisted of six runs. The experiment was designed with the Psychopy 2022.2.5 package (Peirce et al., 2019).

To evaluate color display filter conditions, CVD participants underwent two sessions. While the first session utilized original stimuli in all runs, the second session consisted of two runs for each of three conditions: colors modified by the Windows filter, by the optimized filter and original stimuli. We compared neural responses of CVD participants under two filter conditions with the responses of HC participants to the original stimuli.

We acquired the MRI data focusing on the occipital lobe with a Siemens 3T MAGNETOM Cima.X Scanner with BioMatrix 3T coils. To constrain FOV to posterior occipital lobe, we set twenty-four oblique slices perpendicular to calcarine sulcus (Ryu & Lee, 2024). The scanner was set as 1.5s TR, 30ms TE, 75° FA, $2 \times 2 \times 2$ mm voxel size, and 96×80 matrix size.

3 Preprocessing and response estimation

Spatial Normalization and registration

We converted the acquired data into BIDS format, and removed facial features using ezBIDS (Levitas et al., 2024). Functional images were then registered to the T1w image. Coregistration was performed by optimizing a Mutual Information based (MI) cost function (Maes, Collignon, Vandermeulen, Marchal & Suetens, 1997; Wells, Viola, Atsumi, Nakajima & Kikinis, 1996) via FreeSurfer's `mri_coreg`. This method is well suited to images with limited field-of-view (FOV) and obliquity because it does not require complete anatomical boundaries outside the acquired image. Image orientation was first initialized according to the scanner-defined obliquity information in the NIfTI header. The mutual-information cost function was then optimized to maximize the statistical dependence between T1w and BOLD intensity distribution within the overlapping region. T1w-to-MNI normalization was then performed using a 12-degrees-of-freedom affine transformation (FSL FLIRT; Jenkinson, Bannister, Brady & Smith, 2002), followed by nonlinear warping (FNIRT; Andersson, Jenkinson & Smith, 2007), yielding final BOLD data in MNI space at 2 mm isotropic resolution.

ROI definition

Bilateral regions of interest (V1, V2, V3, and hV4; Fig. 1b) were defined in MNI space at 2 mm resolution using the Wang probabilistic atlas (Wang, Mruczek, Arcaro & Kastner, 2015). Bilateral ROIs were thresholded at 50% atlas probability (following Wang et al., 2015) and intersected with each individual's BOLD brain mask. Voxel counts varied across ROIs (V1: 655 ± 214 ; V2: 451 ± 145 ; V3: 103 ± 29 ; hV4: 63 ± 22 ; all mean $\pm$ SD across subjects).

FIR-based response estimation

Neural response amplitudes were estimated using a two-stage procedure: first extracting a hemodynamic response function (HRF) for each ROI via a finite impulse response (FIR) model, followed by a general linear model (GLM; Dale, 1999; Brouwer & Heeger, 2009, 2013). This data-driven approach recovers a more accurate HRF than canonical models because the FIR model makes no assumption about HRF shape. In the first stage, we estimated the HRF of each ROI using an FIR model. Design matrices for each run were constructed with all stimulus onsets regardless of color identity. The FIR basis set spanned 8 TRs (12 s). The FIR beta weights were estimated by fitting the model to the preprocessed voxel time courses using linear regression. A mean ROI-level HRF was then obtained by averaging the FIR beta weights across voxels within each ROI. Voxels in the top 50% of variance explained by this mean ROI-level HRF were retained, balancing signal quality with sample size (mean retained: V1 328 ± 107 , V2 226 ± 73 , V3 52 ± 15 , hV4 32 ± 11 voxels).

In the second stage, a design matrix was constructed by convolving binary time courses for each color condition with the estimated ROI-specific HRF and its temporal derivative. The temporal derivative was included to capture voxel-to-voxel variability in HRF timing. This design matrix was then fitted to the preprocessed MNI-space voxel time courses using linear regression to estimate response amplitudes. Beta weights for the derivative regressors were discarded, and the remaining coefficients provided a single response amplitude estimate for each voxel, color, and run.

Within-subject Procrustes alignment

To improve cross-run consistency in voxel response patterns within each participant, voxel response amplitudes were aligned across runs using Procrustes alignment (Gower, 1975). This preprocessing reduces measurement noise across runs while preserving color representational geometry, thereby improving the quality of subsequent SRM fitting. To preserve structural representations, voxel response amplitude matrices (across all colors) from runs 2–6 were aligned to run 1, selected as reference to maintain temporal ordering consistency, by estimating an orthogonal transformation that minimized the residual error (Equation 1). To preserve voxel response amplitude magnitudes, scaling was prohibited. To evaluate whether alignment improved cross-run consistency while preserving stimulus-related structure, we computed the mean pairwise Pearson correlation of voxel response amplitude patterns across all run pairs (15 comparisons). Average correlation improved from $r = 0.54$ (SD = 0.12) before alignment to $r = 0.71$

(SD = 0.09) after alignment, indicating improved cross-run consistency.

$$X_{\text{aligned}} = XQ \quad (1)$$

where $X \in \mathbb{R}^{n_{\text{voxels}} \times n_{\text{colors}}}$ is the amplitude matrix for a given run, $X_{\text{ref}} \in \mathbb{R}^{n_{\text{voxels}} \times n_{\text{colors}}}$ is the reference matrix, and $Q \in \mathbb{R}^{n_{\text{voxels}} \times n_{\text{voxels}}}$ is the orthogonal matrix minimizing $\|X_{\text{ref}} - XQ\|_F$.

4 Functional Alignment and shared representational space

Construction of HC-derived common space with Shared Response Model

After within-subject Procrustes alignment, cross-subject alignment was performed using shared response modeling (SRM; Chen et al., 2015), a dimension-reduction technique related to hyperalignment (Haxby et al., 2011; Guntupalli et al., 2016). This enabled comparison of color representations in a common space across participants with different numbers of voxels. SRM learns a low-dimensional common space S where between-subject variability is minimized. We utilized the Brain Imaging Analysis Kit package (BrainIAK) for SRM implementation.

SRM estimates linear mappings between each participant’s original voxel space and the common space. For participant i , let $X_i \in \mathbb{R}^{v \times c}$ denote the matrix of fMRI response amplitudes, where v is the number of voxels in the ROI and $c = 8$ colors. SRM estimates a common space $S \in \mathbb{R}^{k \times c}$ and participant-specific mappings relating each X_i to that space, where k is the reduced dimension.

$$X_i = W_i S + E_i \quad (2)$$

$$\min_{W_i, S} \sum_{i=1}^N \|X_i - W_i S\|_F^2 \quad \text{s.t.} \quad W_i^T W_i = I_k \quad (3)$$

In equation 3, $W_i \in \mathbb{R}^{v \times k}$ is the orthogonal mapping matrix mapping participant i ’s voxel space to the common space, and E_i is the residual error. The matrix W_i is optimized to minimize the sum of Frobenius norm of residuals across all participants. The orthogonal constraint ($W_i^T W_i = I_k$) helps preserve geometric relationships between colors—distances and angles—in the common space. To establish an HC-derived reference common space, we fitted SRM using voxel response data from the seven HC participants only. CVD participants were subsequently projected into this fixed space by computing $W_{\text{CVD}} = UV^T$,

where $U\Sigma V^T = \text{SVD}(X_{\text{CVD}} \cdot S^\dagger)$, yielding an orthonormal mapping that best aligned CVD data to the HC-derived space without re-estimating S .

To identify which color pairs showed geometric deviations in CVD, we computed the representational dissimilarity matrix (RDM) for each of all 28 color pairs in the common space using correlation distance (1 - Pearson r), which captures pattern similarity while being robust to amplitude scaling and additive offsets. To identify pairs exhibiting significant deviation, we computed element-wise differences between each CVD participant's RDM and the mean HC RDM, evaluating significance via 95% bootstrap confidence intervals (1,000 iterations, resampling HC participants with replacement). The reduced dimension k was set to: V1 $k = 4$; V2 $k = 4$; V3 $k = 3$; hV4 $k = 3$.

5 Color decoding and voxel response prediction model

Forward Encoding Model

We modeled hue-selective neural populations using half-wave rectified and squared sinusoidal basis functions following Brouwer and Heeger (2009) and Brouwer and Heeger (2013) (**Fig 2a**). Each of six hypothetical channels was tuned to a preferred hue θ_k evenly spaced around the hue circle, with response profile $r_k(\theta) = \max(0, \cos(\theta - \theta_k))^2$. The squaring operation sharpens the tuning profiles and introduces higher-order harmonic components, requiring multiple phase-shifted channels to fully span the circular hue space. Six channels provide a complete basis set, allowing any hue-selective tuning curve to be expressed as a weighted sum of the basis functions (Freeman & Adelson, 1991).

This forward model formalizes how hypothetical color channels map to observed voxel responses: each voxel's response B is modeled as a weighted sum of six channel outputs C , where the weight matrix W represents the brain's mapping from hue-selective population activity to measured BOLD signals (Equation 4).

$$B = WC \quad (4)$$

Training and evaluation were conducted using cross-validations. Preprocessed response data, represented either in voxel space or in the SRM-derived shared-space, were partitioned into training (B_1, C_1) and test sets (B_2, C_2). During training, the weight matrix W was estimated via ridge regression (Equation 5), where the regularization parameter α was selected via generalized cross-validation (GCV; Golub, Heath & Wahba, 1979). Once W was estimated, channel outputs were decoded from test voxel responses (Equation 6a), or voxel responses were predicted from channel outputs

(Equation 6b). Decoded channel outputs $\hat{C}$ were converted to hue angle estimates by finding the angle θ that minimized $\|\hat{C} - C(\theta)\|_2$.

$$\hat{W} = B_1 C_1^T (C_1 C_1^T + \alpha I)^{-1} \quad (5)$$

$$\text{(color decoding)} \quad \hat{C} = (\hat{W}^T \hat{W})^{-1} \hat{W}^T B_2 \quad (6)$$

$$\text{(voxel prediction)} \quad \hat{B} = \hat{W} C_2 \quad (7)$$

Cross-validation: discrimination vs interpolation

To dissociate color discrimination from continuous hue-space interpolation, we evaluated the forward model using two complementary cross-validation schemes.

Leave-one-run-out (LORO) tested whether individual colors could be reliably distinguished from voxel responses, assessing discrimination consistency across stimulus repetitions. For each of six runs, the corresponding response matrices (B, C) were held out as the test set, W was trained on the remaining five runs, and voxel responses were predicted (Equation 6b).

Leave-one-color-out (LOCO) tested whether the model could interpolate to novel hues not seen during training, assessing the preservation of continuous color-space geometry. For each color, all six run response matrices of the target color were held out. W was trained on the remaining seven colors, and voxel responses to the held-out color were predicted (Equation 6b).

Both were evaluated using Pearson correlation between predicted and observed response vectors, computed across all voxels and test samples. LORO performance indicates color discriminability, while LOCO performance reflects the intactness of continuous hue-space geometry.

6 Behavioral-neural concordance

Behavioral tasks

To examine the relationship between geometrical deviations in neural representations and perceptual color discrimination, participants completed two subsequent behavioral tasks under display settings matching the fMRI environment.

First, participants performed a same-different color discrimination task to estimate their Just Noticeable Difference (JND) thresholds. On each trial, participants saw a pair of colors and reported whether they were identical or different. Stimuli were transformed from CIEL*a*b* space into polar coordinates to allow continuous interpolation along the hue angle. Thresholds were estimated using two interleaved 1-up/1-down adaptive staircases per color pair, terminating after eight

reversals (Levitt, 1971). The task included eight color pairs selected based on RDM deviation analysis results: (1) three pairs showing significant difference in both deuteranopes and protanomalous compared to HC participants (Yellow-purple, Blue-purple, red-orange), (2) two pairs that significantly differed only in deuteranopes (Orange-yellow, Yellow-green), (3) two pairs that significantly differed only in protanomalous (Cyan-magenta, Green-blue), and (4) a control pair that did not differ significantly across any group (Red-cyan).

Second, participants performed an 8-alternative forced-choice (8-AFC) color identification task. On each trial, a circular color patch was presented, and participants selected the matching hue from eight options identical to the fMRI stimulus set.

7 Filter design

We modeled CVD color distortion as a cone-specific spectral shift ($\Delta\lambda$ in nm) applied to the stimulus hue angles (Brettel et al., 1997). The filter parameters were optimized to minimize the discrepancy between CVD neural responses to filtered stimuli and HC responses to original stimuli. Validation was performed using both RDM convergence and V4 LOCO performance as dual criteria.

8 Reproducibility

All analyses were conducted in Python 3.10 with numpy 1.24.3, scipy 1.11.3, scikit-learn 1.3.0, and BrainIAK 0.11. Random seeds were fixed (seed=42) for permutation tests and bootstrap resampling. Analysis code and preprocessing scripts are available at <https://github.com/haba6030/colorBlindanalysis>. Anonymized preprocessed data are available upon reasonable request to the corresponding author.

9 Results

9.1 Behavioral Performance

All participants successfully completed the attention task during fMRI scanning (mean accuracy: HC = XX%, CVD = XX%). Response times did not differ significantly between groups ($t = XX$, $p = XX$), indicating comparable task engagement and alertness throughout the scanning session.

In the post-scan color identification task, HC participants correctly identified XX% of colors on average, while CVD participants showed reduced accuracy (XX%), particularly for colors along the red-green axis. Just noticeable difference (JND) thresholds, estimated using adaptive staircases (Levitt, 1971), revealed elevated thresholds for CVD participants on specific color pairs identified from the neural RDM analysis (see be-

low).

9.2 Neural Color Representations in Shared Space

We compared neural color representations between HC and CVD participants using shared response modeling (Chen et al., 2015), which projects individual voxel patterns into a low-dimensional common space. Following the approach of haxby2011 and guntupalli2016, we estimated a common space from HC participants only and subsequently projected CVD data into this HC-derived space.

Individual CVD participants showed increased disparity from the HC-derived common space compared to leave-one-out HC controls. Using crawford1998 modified t -tests appropriate for single-case comparisons, we found significant deviations in early visual cortex: V1 (CVD-1: $t(6) = XX$, $p = XX$; CVD-2: $t(6) = XX$, $p = XX$) and V2 (CVD-1: $t(6) = XX$, $p = XX$; CVD-2: $t(6) = XX$, $p = XX$).

Representational dissimilarity matrix (RDM) analysis, quantified using correlation distance, revealed specific color pairs that differed significantly between CVD and HC groups. Significance was assessed via 95% bootstrap confidence intervals (1,000 iterations), resampling HC participants with replacement to account for small sample variability. In V1, CVD participants showed altered dissimilarity for [specific color pairs determined by empirical analysis]. Similar patterns emerged in V2 and hV4, with effect sizes largest in early visual areas.

V1: [Insert specific results with effect sizes]

V2: [Insert specific results with effect sizes]

V3: [Insert specific results with effect sizes]

hV4: [Insert specific results with effect sizes]

9.3 Forward Encoding Model Performance

We evaluated forward encoding models using three complementary cross-validation schemes, following the approach of brouwer2009 and brouwer2013. Models were estimated via ridge regression with the regularization parameter selected by generalized cross-validation (Golub et al., 1979).

9.3.1 Color Discrimination: Leave-One-Run-Out (LORO)

Leave-one-run-out cross-validation demonstrated that individual colors could be successfully decoded from voxel responses in all ROIs. Decoding accuracy exceeded chance (12.5%) for both HC (V1: XX%, V2: XX%, V3: XX%, hV4: XX%) and CVD participants (V1: XX%, V2: XX%, V3: XX%, hV4: XX%). Importantly, group comparisons via nonparametric permutation tests (Nichols & Holmes, 2002) revealed no

significant differences in LORO performance ($p < .05$ for all ROIs), indicating that color discrimination performance was preserved in CVD.

9.3.2 Color Interpolation: Leave-One-Color-Out (LOCO)

Leave-one-color-out cross-validation tested whether forward encoding models could interpolate to novel, held-out colors—a critical test of whether the geometric structure of neural color space is preserved. HC participants showed successful interpolation in higher-order visual areas (hV4: $r = .88$, $p < .001$), replicating prior findings (Brouwer & Heeger, 2009). However, CVD participants showed significantly reduced LOCO performance compared to HC in early visual areas: V1 (Hedges' $g = .45$, $p = .001$) and V2 (Hedges' $g = .45$, $p = .001$). This dissociation—preserved discrimination (LORO) but impaired interpolation (LOCO)—suggests that CVD distorts the continuous structure of neural color space rather than simply reducing color discriminability.

9.3.3 Zero-Shot Generalization: Leave-One-Subject-Out (LOSO)

Leave-one-subject-out analysis tested whether the HC group-derived common space (Chen et al., 2015) supported zero-shot generalization to novel individuals. Consistent with prior hyperalignment studies (Haxby et al., 2011; Guntupalli et al., 2016), LOSO performance was above chance in hV4 (HC: $r = .88$; CVD: $r = .88$). Critically, there was no significant difference between HC and CVD in LOSO performance ($p = .99$), indicating that the shared representational space was equally applicable to both groups. This finding validates our use of an HC-derived common space for CVD comparisons and suggests that the fundamental structure of neural color coding is shared across individuals despite peripheral differences.

9.4 Cone-Shift Model Validation

[To be completed based on Phase 2 cone-shift pipeline results]

Using a cone-specific spectral shift model (Brettel et al., 1997), we tested whether CVD-related neural distortions could be predicted by simulating altered cone spectral sensitivities. [Add specific results here after analysis completion.]

10 Discussion

10.1 Summary of Main Findings

The present study investigated neural color representations in individuals with color vision deficiency (CVD)

using multivariate pattern analysis and computational modeling. Our key findings were threefold. First, CVD participants showed systematic deviations from the HC-derived common representational space, particularly in early visual cortex (V1, V2). Second, while color discrimination (LORO) was preserved in CVD, color interpolation (LOCO) was impaired in early visual areas. Third, the cone-shift model successfully predicted and characterized these neural distortions.

10.2 Neural Representations in CVD: Comparison with Previous Literature

Our finding of altered neural color representations in CVD extends previous behavioral and computational modeling studies by demonstrating that perceptual differences have measurable neural correlates throughout the visual hierarchy. The pattern of results—preserved discrimination but impaired interpolation—suggests that CVD affects the geometric structure of color space rather than simply reducing the discriminability of individual colors. This finding is consistent with computational models of CVD that emphasize distortions in color space rather than mere confusion between colors (Brettel et al., 1997).

These findings build directly on the seminal work of Brouwer (2009), who demonstrated successful novel-color reconstruction in V4/VO1 using forward encoding models in individuals with typical color vision. Our results extend this framework by showing that interpolation capability is selectively impaired in CVD despite preserved classification performance. Critically, Brouwer (2009) did find that V1 showed weaker interpolation compared to V4 even in typical observers, consistent with our finding that early visual areas may encode local color discriminability without fully representing the continuous structure of color space.

More recent work by Brouwer (2013) demonstrated categorical clustering of neural color representations in ventral temporal cortex. Our findings are complementary: while categorical boundaries may be preserved in CVD (as evidenced by intact discrimination), the continuous metric structure between categories appears distorted. This dissociation between categorical and continuous aspects of color coding deserves further investigation in future work.

10.3 Dissociation Between Geometric and Functional Distortions

Importantly, we observed a dissociation between SRM-based geometric distortion (which was modest) and LOCO-based functional impairment (which was pronounced). This finding has important methodological implications. Representational dissimilarity matrices (RDMs), while widely used to character-

ize neural representational geometry (Haxby et al., 2011), may not fully capture functional consequences of representational distortions. In our data, only LOCO performance—a direct measure of interpolation capacity—predicted behavioral JND thresholds. This suggests that RDMs reflect metric properties of neural space, while cross-validated model performance reflects functional interpolation capacity. Future studies should carefully distinguish between these complementary measures.

The hyperalignment framework developed by haxby2011 and extended by guntupalli2016 provides a powerful approach for comparing representations across individuals. Our finding that LOCO performance did not differ between HC and CVD validates the use of HC-derived common spaces for CVD comparisons and suggests that the fundamental structure of neural color coding is shared across individuals despite peripheral differences in cone spectral sensitivities.

10.4 Hierarchical Organization of Color Interpolation

The dissociation between LORO (discrimination) and LOCO (interpolation) performance is theoretically significant. It demonstrates that neural populations in V1/V2 can discriminate colors they cannot interpolate—suggesting that these regions encode local color contrasts but may not represent the full geometric structure of perceptual color space. In contrast, hV4 showed successful interpolation, consistent with proposals that this region serves as a perceptual hub for color processing. This hierarchical organization mirrors findings in other domains, such as object recognition, where early visual areas encode local features while higher areas encode more abstract, generalizable representations (Guntupalli et al., 2016).

Our analysis of different basis functions revealed a tradeoff between dimensionality and interpolation performance. V1/V2 showed negative correlations between k (number of channels) and LOCO performance, while hV4 showed positive correlations. This pattern suggests that optimal channel configurations differ across the visual hierarchy—consistent with proposals that early areas encode color with high resolution but limited generalization, while higher areas trade resolution for generalizability.

10.5 Cone-Shift Model and Clinical Implications

Our cone-shift model provides a principled framework for designing personalized color correction filters. Unlike previous approaches that rely solely on perceptual matching or heuristic transformations, our method uses neural data to optimize filter parameters, potentially

enabling more effective correction of CVD-related distortions. The finding that LOCO performance predicted behavioral JND thresholds (100% concordance) validates this approach and suggests that neural interpolation capacity is a useful biomarker for perceptual discrimination ability.

However, an important caveat is that our cone-shift model was validated only on neural data, not in behavioral experiments with actual corrective filters. Future work must test whether filters optimized to restore neural interpolation actually improve perceptual discrimination in real-world settings. This validation is critical before clinical applications can be pursued.

10.6 Limitations and Future Directions

The primary limitation of this study is the small CVD sample size ($n = 2$). While we used single-case statistics (Crawford & Howell, 1998) to rigorously test individual effects, replication in a larger sample is needed to establish population-level conclusions. Future studies should recruit larger samples of both deuteranopes and protanomalous individuals to test whether CVD subtype modulates neural representations. Our preliminary data suggest differential patterns (e.g., sub-08 deutan vs. sub-09 protan showed different color pair distortions), but more participants are needed to draw firm conclusions.

We focused on half-wave rectified sinusoidal basis functions following brouwer2009 and brouwer2013. While we compared multiple basis configurations, non-linear models or data-driven approaches (e.g., deep neural networks) may capture additional structure in neural color representations. However, our permutation analyses suggest that linear models provide a good first-order approximation, and the theoretical interpretability of basis function models remains a significant advantage for hypothesis testing.

Although we used CIE $L^*a^*b^*$ space to define stimuli, this space is designed for typical color vision and may not optimally describe CVD color perception. Future work could investigate whether CVD-specific perceptual spaces (e.g., based on multidimensional scaling of behavioral judgments) better align with neural representations. However, using a standard color space facilitates comparison with prior literature and enables principled stimulus design.

Finally, while our cone-shift model showed promising results in predicting neural distortions, validation of designed filters in behavioral and neural experiments is an important next step. This would test whether filters optimized to restore neural interpolation also improve perceptual discrimination and whether neural changes predict behavioral improvements.

10.7 Conclusions

Color vision deficiency alters neural color representations in a systematic and hierarchically organized manner. While color discrimination is preserved, the geometric structure of color space is distorted, particularly in early visual cortex. Forward encoding models provide a powerful framework for characterizing these distortions and may enable personalized approaches to color correction. These findings advance our understanding of how genetic variation in photoreceptor spectral sensitivity shapes cortical color representations and demonstrate that multivariate fMRI analysis can reveal functional consequences of sensory deficits that are not apparent in simple activation measures.

References

- Andersson, J. L., Jenkinson, M. & Smith, S. (2007). *Non-linear registration aka spatial normalisation* (Technical Report No. TR07JA2). FMRIB.
- Brettel, H., Viénot, F. & Mollon, J. D. (1997). Computerized simulation of color appearance for dichromats. *Journal of the Optical Society of America A*, 14(10), 2647–2655. doi: 10.1364/JOSAA.14.002647
- Brouwer, G. J. & Heeger, D. J. (2009). Decoding and reconstructing color from responses in human visual cortex. *Journal of Neuroscience*, 29(44), 13992–14003. doi: 10.1523/JNEUROSCI.3577-09.2009
- Brouwer, G. J. & Heeger, D. J. (2013). Categorical clustering of the neural representation of color. *Journal of Neuroscience*, 33(39), 15454–15465. doi: 10.1523/JNEUROSCI.2472-13.2013
- Chen, P. H., Chen, J., Yeshurun, Y., Hasson, U., Haxby, J. & Ramadge, P. J. (2015). A reduced-dimension fmri shared response model. In *Advances in neural information processing systems* (Vol. 28, pp. 460–468).
- Commission Internationale de L'Eclairage. (1986). *Colorimetry* (2nd ed.). CIE. (Publication No. 15.2)
- Crawford, J. R. & Howell, D. C. (1998). Comparing an individual's test score against norms derived from small samples. *The Clinical Neuropsychologist*, 12(4), 482–486. doi: 10.1076/clin.12.4.482.7241
- Dale, A. M. (1999). Optimal experimental design for event-related fmri. *Human Brain Mapping*, 8(2-3), 109–114. doi: 10.1002/(SICI)1097-0193(1999)8:2/3<109::AID-HBM7>3.0.CO;2-W
- Durnez, J., Moerkerke, B. & Nichols, T. E. (2018). neurodesign: optimal experimental designs for task fmri. *Journal of Open Source Software*, 3(32), 564. doi: 10.21105/joss.00564
- Freeman, W. T. & Adelson, E. H. (1991). The design and use of steerable filters. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 13(9), 891–906. doi: 10.1109/34.93808
- Golub, G. H., Heath, M. & Wahba, G. (1979). Generalized cross-validation as a method for choosing a good ridge parameter. *Technometrics*, 21(2), 215–223. doi: 10.1080/00401706.1979.10489751
- Gower, J. C. (1975). Generalized procrustes analysis. *Psychometrika*, 40(1), 33–51. doi: 10.1007/BF02291478
- Guntupalli, J. S., Hanke, M., Halchenko, Y. O., Connolly, A. C., Ramadge, P. J. & Haxby, J. V. (2016). A model of representational spaces in human cortex. *Cerebral Cortex*, 26(6), 2919–2934. doi: 10.1093/cercor/bhw068
- Haxby, J. V., Guntupalli, J. S., Connolly, A. C., Halchenko, Y. O., Conroy, B. R., Gobbini, M. I., ... Ramadge, P. J. (2011). A common, high-dimensional model of the representational space in human ventral temporal cortex. *Neuron*, 72(2), 404–416. doi: 10.1016/j.neuron.2011.08.026
- Ishihara, S. (1917). *Tests for colour-blindness*. Tokyo: Kanehara Shuppan.
- Jenkinson, M., Bannister, P., Brady, M. & Smith, S. (2002). Improved optimization for the robust and accurate linear registration and motion correction of brain images. *NeuroImage*, 17(2), 825–841. doi: 10.1016/s1053-8119(02)91132-8
- Levitas, D., Hayashi, S., Vinci-Booher, S., Heinsfeld, A., Bhatia, D., Lee, N., ... Pestilli, F. (2024). ez bids: Guided standardization of neuroimaging data interoperable with major data archives and platforms. *Scientific Data*, 11, 179. doi: 10.1038/s41597-024-02959-0
- Levitt, H. (1971). Transformed up-down methods in psychoacoustics. *The Journal of the Acoustical Society of America*, 49(2), 467–477. doi: 10.1121/1.1912375
- Maes, F., Collignon, A., Vandermeulen, D., Marchal, G. & Suetens, P. (1997). Multimodality image registration by maximization of mutual information. *IEEE Transactions on Medical Imaging*, 16(2), 187–198. doi: 10.1109/42.563664
- Nichols, T. E. & Holmes, A. P. (2002). Nonparametric permutation tests for functional neuroimaging: a primer with examples. *Human Brain Mapping*, 15(1), 1–25. doi: 10.1002/hbm.1058
- Peirce, J., Gray, J. R., Simpson, S., MacAskill, M., Höchenberger, R., Sogo, H., ... Lindeløv, J. K. (2019). Psychopy2: Experiments in behavior made easy. *Behavior Research Methods*, 51(1),

- 195–203. doi: 10.3758/s13428-018-01193-y
- Ryu, J. & Lee, S.-H. (2024). Slice planning perpendicular to calcarine sulcus for retinotopic mapping. *Journal of Neuroscience Methods*. (In preparation)
- Wang, L., Mruczek, R. E., Arcaro, M. J. & Kastner, S. (2015). Probabilistic maps of visual topography in human cortex. *Cerebral Cortex*, 25(10), 3911–3931. doi: 10.1093/cercor/bhu277
- Wells, W. M., Viola, P. A., Atsumi, H., Nakajima, S. & Kikinis, R. (1996). Multi-modal volume registration by maximization of mutual information. *Medical Image Analysis*, 1(1), 35–51. doi: 10.1016/S1361-8415(01)80004-9