

Preserved Discrimination but Impaired Interpolation in hV4 Identifies the Correctable
Distortion in Colour Vision Deficiency

Your Name

Your University

Abstract

Placeholder abstract. Replace this with your actual abstract text. This study investigated neural color representations in individuals with color vision deficiency (CVD) using functional MRI and computational modeling approaches.

Keywords: color vision deficiency, fMRI, shared response model, forward encoding model, representational similarity analysis

Preserved Discrimination but Impaired Interpolation in hV4 Identifies the Correctable Distortion in Colour Vision Deficiency

Introduction

Color vision deficiency (CVD) affects approximately 8% of males and 0.5% of females of Northern European descent (Bosten, 2019) and arises from altered spectral sensitivity of the L- or M-cone photopigment, driven by X-linked opsin-gene polymorphisms (Neitz & Neitz, 2011; Deeb, 2005). In anomalous trichromats the deviant cone’s peak sensitivity is displaced by approximately 2–12 nm from its normal value (Stockman & Sharpe, 2000). This shift reduces the L–M cone-opponent signal rather than abolishing it. The peripheral consequences — reduced red–green discrimination, category confusions, and elevated just-noticeable differences along the L–M axis — are the classical phenotype captured by the Ishihara test (Ishihara, 1917) and by subsequent psychophysics (Bosten, 2019).

CVD is not an on/off loss: downstream cortical color machinery remains intact and receives a quantitatively reweighted but qualitatively similar input. This makes CVD a paradigmatic example of a condition in which a high-dimensional neural representation is not absent but *structurally distorted*. The scientific and translational question is therefore not whether color information survives the retina — it manifestly does — but whether, and how, its cortical geometry is reshaped at the level of the individual. This question remains open, and an individualized assistive filter must answer it before it can be designed.

CVD correction filters in current use fall into two structurally related families. Hardware filters — dichroic notch lenses such as the EnChroma line — attenuate spectral bands near the L–M crossover to widen the post-receptoral L–M signal. Software filters operate at the display: Brettel–Viénot–Mollon simulation (Brettel, Viénot & Mollon, 1997) and the parameterized Machado retinal model (Machado, Oliveira & Fernandes, 2009) forward-project a stimulus onto an anomalous retina. Their Daltonization inverses, including recent learned reformulations (Akalin et al., 2025), redistribute the lost contrast into channels the user can still see. Despite this

diversity, the family shares two structural commitments: optimization against a *population-average* retinal model, with the same transformation applied to every user.

Attempts to escape this average-user ceiling have predominantly taken one route: tuning a retinal-model parameter — the simulated L- or M-cone shift in Machado-class models — to a user’s responses on plate-style or appearance-matching tests. The route is constrained at its target: its optimization criterion is the user’s appearance *report*, not the cortical representation that downstream tasks read out.

This distinction matters because long-term cortical adaptation in anomalous trichromats reshapes hue appearance independently of the retina (Tregillus et al., 2021). Appearance-matched calibration therefore inherits whatever compensation the cortex has already imposed. A device tuned to undo only the retinal component then re-renders the input *as if* that cortical compensation did not exist. Consistent with this account, clinical validation of consumer color-enhancing CVD-correction glasses reports at most modest changes in chromatic discrimination thresholds in congenital red–green deficiencies (Patterson et al., 2022), and recent behavioral work likewise shows that population-average filters change the appearance of color without substantively shifting discrimination thresholds at threshold contrast (Somers, Franklin & Bosten, 2024), while the underlying CVD phenotype varies markedly even within a single diagnostic category (Bosten, 2019). What is missing from this landscape is a filter derived from, and validated on, the *user’s own* neural color geometry.

Cortical color representation is a distributed, hierarchical computation (Gegenfurtner, 2003; Shapley & Hawken, 2011; Conway, Eskew, Martin & Stockman, 2018). Early visual cortex carries cone-opponent signals. Higher-order regions, in particular hV4 (Brouwer & Heeger, 2009; Bannert & Bartels, 2018), support a population code for continuous hue that can be read out by forward encoding models (Brouwer & Heeger, 2009; Parkes, Marsman, Oxley, Goulermas & Wuerger, 2009; Kuriki, Sun, Ueno, Tanaka & Cheng, 2015). Cross-subject shared-response modelling has recently been validated for healthy color decoding (Bannert & Bartels, 2025), providing a methodological scaffold that the present clinical extension adopts.

Behavioral work converges on the conclusion that anomalous trichromats do not merely discriminate less well; they compensate. Long-term cortical adaptation reshapes hue appearance (Tregillus et al., 2021; Boehm, MacLeod & Bosten, 2014; Webster, 2015), and hue-scaling experiments in anomalous trichromats ($n = 10$ AT vs. 26 NT) reveal a consistent 21.4° rotation of the blue–yellow response peak — a phase shift of the half-rectified cosine fit to scaling responses across the 360° stimulus circle (NT peak 127.5° vs. AT peak 106.1° ; $t(34) = 5.95$, $p < 0.001$) (Emery, Volbrecht, Peterzell & Webster, 2021; Isherwood, Joyce, Parthasarathy & Webster, 2020). These results locate the key variance not at the retina alone but at the retinal–cortical interface.

Yet three concrete gaps remain. (*Gap 1*) Existing neural studies of color characterize population-mean geometry in healthy observers (Brouwer & Heeger, 2009; Bannert & Bartels, 2018; Kuriki et al., 2015) or, in CVD, group means at the level of activation magnitude. To our knowledge — based on PubMed and Google Scholar searches combining “CVD” / “color vision deficiency” / “anomalous trichromacy” with “fMRI”, “forward encoding”, and “individual differences” — none has inverted the high-dimensional cortical signal into an interpretable, individualized distortion field. (*Gap 2*) CVD neural work has relied principally on classification accuracy, which is a sufficient statistic for category identity but discards the continuous relational structure (Brouwer & Heeger, 2013) that downstream perceptual and translational tasks actually demand. (*Gap 3*) Behavioral evidence for cortical compensation has not been translated into a concrete criterion for choosing between retinal-level and cortical-level correction strategies — an omission that directly limits filter design.

A filter for a continuous-hue condition must operate on continuous-hue representation, not on categorical labels. This observation motivates a complementary pair of cross-validation schemes on the same forward encoding model. Leave-one-run-out (LORO) cross-validation probes whether the 8 trained hues can be *classified* from held-out runs. It measures category-level discriminability and is the minimum requirement for a filter to matter: if the target cortex fails to discriminate the displayed colors, no stimulus-space adjustment has leverage on which to act.

Leave-one-color-out (LOCO) cross-validation, by contrast, probes whether the encoder *interpolates* from its 7 training hues to a held-out hue. It measures the integrity of the continuous manifold on which any hue correction must act.

A joint pattern is therefore predicted and empirically required: LORO *preserved* in CVD (the filter precondition) together with LOCO *impaired* (the filter target). Finding only the first would make a filter redundant; finding only the second would preclude any stimulus-space remedy. Finding both together defines the regime in which individualized corrective filters are simultaneously warranted and tractable, and in which the per-hue LOCO vulnerability vector — the subject-specific signature of where the continuous geometry fails — becomes the phenotype to be inverted.

We test whether *individual hV4 color geometry* carries person-specific distortion through a single-case fMRI characterization of two CVD adults (one moderate deutan, one moderate-severe protan) under Crawford–Howell single-case inference (see Methods), alongside a normative reference of seven healthy controls. Findings specific to sub-09 — the single protan participant in our cohort — are reported as exploratory single-case observations requiring independent replication.

1. Does the continuous cortical hue geometry in CVD carry

individually-structured distortion? We use the LORO/LOCO dissociation on hV4 forward encoding to distinguish preserved discrimination from impaired interpolation, and extract a per-hue vulnerability vector $\mathbf{v} \in \mathbb{R}^8$ as each individual’s geometric fingerprint.

2. Can that distortion be expressed in a small number of physiologically interpretable parameters, and does an independent geometric criterion corroborate the fit?

We infer 1–2 DOF parameters from $\mathbf{v}$ under three mechanistic model families (retinal Machado, retinal–cortical R+C, cortical 2-component) with closed testing. We then cross-check the winning fit against two independent signatures. The first is the ΔRDM , the pairwise difference of representational dissimilarity matrices (RDMs) in the common space derived by the Shared Response Model (SRM; Chen et al., 2015). The second is the

cross-cohort S-cone compensation angle described above.

- 3. Does the pre-image of the fitted distortion constitute a testable corrective filter, and do retinal and cortical hypotheses agree on what to display?** We compute the per-hue pre-image of each fit and assess bijectivity. Retinal and cortical fits converge on *who* is distorted but diverge on *how* to correct — generating a falsifiable behavioral prediction to be tested in a same-day 2AFC discrimination experiment.

We hypothesized that (i) categorical hue discrimination at hV4 is preserved in both CVD individuals while continuous-hue interpolation is selectively impaired along S-cone-intermediate hues; (ii) the resulting per-hue vulnerability profile is captured by a small number of physiologically interpretable parameters anchored on the S-cone opponent axis; and (iii) the pre-image of these parameters yields a stimulus-space correction filter whose direction differs between retinal-level and cortical-level model fits. Together, the three steps yield a per-individual filter whose input (cortical signal), intermediate parameters, and output (stimulus shift) are each testable in a single CVD user. Per-subject estimates are validated against an HC leave-one-out distribution under the same loss (Supplementary); case-level distinctness rests on joint anchors — HC parameter-norm bracket, directional consistency with cone-physiology, and (pending) Phase-3 behavioral validation.

Methods

The analysis follows three stages. First, two neural measures characterised each CVD participant’s hue-processing deficit: leave-one-colour-out (LOCO) decoding at hV4 yielded a per-hue vulnerability profile (§), and the between-group RDM difference (Δ RDM) captured the per-pair geometric deviation of CVD representational geometry from the HC mean (§). Second, two candidate distortion model classes were fit by minimising a combination of neural and behavioural losses, with parameters selected by a pre-specified stability procedure (§-§). Third, the fitted distortion was analytically

inverted to yield a per-subject stimulus-space correction filter (§). This filter was then evaluated in a repeat fMRI session and a 2AFC behavioural discrimination test (§).

Participants

Thirteen volunteers were recruited under Institutional Review Board approval (#2510/002-023) and provided written informed consent. Color vision was screened using the 14-plate Ishihara pseudoisochromatic test (Ishihara, 1917). Participants identifying ≥ 9 plates correctly were classified as healthy controls (HC; $n = 7$, 3 female, age 22.7 ± 2.5 years). Two male participants scoring below this threshold were classified as CVD: sub-08 (deuteranomalous, 5/14 plates corrected) and sub-09 (protanomalous, 7/14 plates corrected). Four additional participants were excluded for failing to complete the full scanning protocol. All CVD results are interpreted as single-case demonstrations using the Crawford and Howell (1998) modified t -test; no population-level inference is made.

Stimuli and task

Eight isoluminant chromatic stimuli were drawn from CIE $L^*a^*b^*$ space at equal angular spacing (0° – 315° , 45° steps; $L^* = 75$, chroma = 40) (Commission Internationale de L’Eclairage, 1986), spanning red, orange, yellow, green, cyan, blue, purple, and magenta (Figure 1A). A gray filler ($a^* = b^* = 0$) served as a null condition.

Stimuli were presented in a Rapid Serial Visual Presentation (RSVP) design to sustain attention and fixation. Each stimulus was a uniform circular chromatic patch (a coloured disc, 1.5 s per stimulus) shown in pseudo-random order optimized with Neurodesign (Durnez, Moerkerke & Nichols, 2018). Participants pressed a button upon detecting two consecutive ‘K’s within a rapid serial letter stream presented at the centre of the chromatic disc. Each run contained eight repetitions of all eight chromatic stimuli and the gray filler; inter-stimulus intervals were drawn from $\{3, 4.5, 6\}$ s. Each session comprised six runs (≈ 7 min each). The experiment was coded in PsychoPy 2022.2.5 (Peirce et al., 2019).

All participants completed JND behavioral testing in the same scanning session as the main MRI acquisition. CVD participants additionally completed a second session in which JND was repeated under the per-subject stimulus-space filter, alongside a second fMRI acquisition (see the Behavioral tasks section).

MRI acquisition and preprocessing

Data were acquired on a Siemens 3T MAGNETOM Cima.X scanner with BioMatrix 3T coils. To constrain coverage to posterior occipital cortex we used 24 oblique slices perpendicular to the calcarine sulcus (Brouwer & Heeger, 2009) (TR 1.5 s, TE 30 ms, FA 75°, $2 \times 2 \times 2$ mm, 96×80 matrix).

Data were converted to BIDS format and defaced with ezBIDS (Levitas et al., 2024). Functional images were registered to the T1w image via mutual-information coregistration (FreeSurfer `mri_coreg`; Maes, Collignon, Vandermeulen, Marchal & Suetens, 1997; Wells, Viola, Atsumi, Nakajima & Kikinis, 1996). T1w-to-MNI normalization used 12-DOF affine (FSL FLIRT; Jenkinson, Bannister, Brady & Smith, 2002) followed by nonlinear warping (FNIRT; Andersson, Jenkinson & Smith, 2007), yielding 2 mm isotropic MNI-space BOLD data.

ROI definition and response estimation

Bilateral V1, V2, V3, and hV4 were defined from the Wang probabilistic atlas (Wang, Mruczek, Arcaro & Kastner, 2015) at $> 50\%$ probability, intersected with each participant’s BOLD brain mask (voxel counts: V1 655 ± 214 ; V2 451 ± 145 ; V3 103 ± 29 ; hV4 63 ± 22 ; mean $\pm$ SD).

Single-trial response amplitudes were estimated with a two-stage GLM (Dale, 1999; Brouwer & Heeger, 2009, 2013). First, an ROI-level hemodynamic response function (HRF) was extracted from a finite-impulse-response (FIR) model (8-TR basis set). Voxels in the top 50% of variance explained by the mean ROI HRF were retained (V1 328 ± 107 ; V2 226 ± 73 ; V3 52 ± 15 ; hV4 32 ± 11 voxels). Second, the estimated HRF was convolved with per-color design matrices and fitted by ridge regression to yield one amplitude per (voxel, color, run), organized as a $6 \times 8 \times V$ array per subject.

Voxel amplitudes were aligned across runs using within-subject Procrustes rotation (Gower, 1975): runs 2–6 were mapped to run 1 via the orthogonal transformation minimizing $\|X_{\text{ref}} - XQ\|_F$ (scaling prohibited).

Shared Response Model

Cross-subject alignment and dimensionality reduction were performed with the Shared Response Model (SRM; Chen et al., 2015). SRM estimates per-participant orthonormal mappings $W_i \in \mathbb{R}^{V \times k}$ (V voxels, k latent dimensions) and a group-shared response matrix $S \in \mathbb{R}^{k \times 8}$:

$$\min_{W_i, S} \sum_{i=1}^N \|X_i - W_i S\|_F^2 \quad \text{s.t.} \quad W_i^\top W_i = I_k \quad (1)$$

SRM was trained exclusively on the seven HC participants to avoid circular inference in CVD comparisons. Each CVD participant was subsequently projected into the fixed HC-derived space by solving $W_{\text{CVD}} = UV^\top$ where $U\Sigma V^\top = \text{SVD}(X_{\text{CVD}} \cdot S^\dagger)$. Reduced dimensions were chosen by mean-rank aggregation across seven-fold leave-one-subject-out (LOSO) cross-validation, using three SRM-quality metrics (RDM split-half reliability, cross-subject RDM correlation, and reconstruction error): V1 $k = 4$; V2 $k = 4$; V3 $k = 3$; hV4 $k = 3$.

Forward encoding model

Both discrimination (LORO) and interpolation (LOCO) decoding used the same forward encoding model (Brouwer & Heeger, 2009). The model represents each stimulus as a $F = 6$ -dimensional channel vector $\mathbf{c}(\theta) = [c_1(\theta), \dots, c_F(\theta)]^\top$, where $c_k(\theta) = \max(0, \cos(\theta - \theta_k))^2$ is a half-wave rectified squared cosine tuned to preferred hue θ_k (preferred hues spaced equally at 60° ; Brouwer & Heeger, 2009). A weight matrix $W \in \mathbb{R}^{F \times V}$ is estimated by ridge regression with the regularization parameter α selected by generalized cross-validation (GCV; Golub, Heath & Wahba, 1979; see Supplementary for implementation details) to predict voxel amplitudes from channel activations: $\hat{X} = CW^\top$. Decoding inverts the encoding direction: for a held-out voxel

pattern $\mathbf{x}$, channel activations are reconstructed as $\hat{\mathbf{c}} = W\mathbf{x}$, and the candidate hue whose basis vector has the highest Pearson correlation with $\hat{\mathbf{c}}$ is returned as the decoded color (Figure 2). Alternative decoders evaluated during model selection are described in Appendix .

Evaluating color discrimination with cross-run decoding

Categorical discrimination was assessed with leave-one-run-out (LORO) cross-validation in SRM-aligned space using the forward encoding model (see the Forward encoding model section). For each of the six runs, W was estimated on the remaining five runs using mean-subtracted SRM projections, and each held-out run’s 8 color patterns were decoded by nearest-neighbour matching. Primary discrimination performance was the 8-class exact accuracy (chance = $1/8 = 0.125$), averaged across six folds.

Group-level HC–CVD comparison used a Mann–Whitney U test on cross-subject generalization scores: the encoding model trained on each HC participant’s SRM-aligned data was applied to every other participant’s aligned patterns, yielding 21 unique HC-to-HC and 14 HC-to-CVD pairwise generalization scores. Within-ROI single-case tests used the Crawford and Howell (1998) modified t -test (one-tailed, lower).

Evaluating color interpolation with cross-color decoding

Continuous hue interpolation was assessed with leave-one-color-out (LOCO) cross-validation using the same forward encoding model (see the Forward encoding model section). Each of the 8 hues was held out in turn; W was estimated on the remaining 7 hues using all 6 runs (pooled 42-sample design matrix), and the held-out hue was predicted by nearest-neighbour matching in channel-output space.

Adjacent accuracy. The primary cross-color decoding metric was adjacent accuracy: the proportion of predictions that fell within ± 1 hue step (45°) of the target hue, averaged across the 6 runs of each test color (range 0–1; chance level under a uniform random predictor = $3/8 = 0.375$; exact-accuracy chance = $1/8 = 0.125$). For HC above-chance testing and for pairwise HC–CVD comparisons at hV4, p -values were

obtained from 1,000 random color-label permutations and from the Crawford and Howell (1998) modified t -test (one-tailed, lower), respectively.

Vulnerability profile. The 8-dimensional vector of per-hue adjacent-accuracy values, $\mathbf{v} \in [0, 1]^8$, constitutes each participant’s *hue vulnerability profile*. For group-level model fit, we report Spearman ρ between the distortion-model-predicted and observed $\mathbf{v}$ (see the Inverse fitting section). Per-hue group comparisons (Figure 3C) used the Crawford and Howell (1998) modified t -test (one-tailed, uncorrected). Effect sizes are reported as Hedges’ d (Hedges & Olkin, 1985) with the Crawford–Howell correction for small- n comparisons.

Between-group representational dissimilarity

To characterize the representational geometry of LOCO impairment in SRM-aligned space, we computed two measures. Mean pairwise disparity captures whether CVD geometry is globally more dispersed than HC. The between-group RDM difference, ΔRDM , captures which color pairs drive the deviation and tests whether a cortical distortion model accounts for the pattern.

Pairwise disparity. For each participant and ROI, the RDM was the 8×8 matrix of pairwise correlation distances ($1 - \text{Pearson } r$). Mean pairwise disparity was the average of the 28 upper-triangle elements. To provide an unbiased HC reference for the held-out HC participant, the HC disparity distribution was estimated by a leave-one-subject-out (LOO) procedure. Under the directional prediction that CVD geometry is more dispersed than HC, individual CVD disparity was compared to the HC LOO distribution by the Crawford and Howell (1998) modified t -test (one-tailed, upper). Significance was assessed at $\alpha = 0.05$.

Geometric deviation from HC and model correspondence. The per-pair deviation of each CVD participant’s RDM from the HC mean was $\Delta\text{RDM} = \text{RDM}_{\text{CVD}} - \overline{\text{RDM}}_{\text{HC}}$, where $\overline{\text{RDM}}_{\text{HC}}$ is the grand mean RDM over the seven HC participants. Correspondence between the observed ΔRDM and the prediction of the two-component distortion model (see Inverse fitting) was reported descriptively as the Spearman ρ between the upper-triangle elements of the two matrices.

Candidate distortion models

Two candidate model classes were evaluated. Each parameterises the angular distortion $\delta\theta(\theta)$ that maps the displayed hue angle θ to the angle perceived by a CVD observer.

Candidate Model 1: retinal-plus-gain (R+C).. The prevailing account attributes CVD colour distortion to a cone-spectral shift at the retinal level (Machado et al., 2009), optionally augmented by a cortical compensation gain g (Boehm et al., 2014; Tregillus et al., 2021):

$$\delta\theta_{\text{RC}}(\theta; g) = (2 - g) \delta\theta_{\text{Mach}}(\theta) \quad (2)$$

where $\delta\theta_{\text{Mach}}$ is the per-hue angular shift predicted by Machado simulation applied to Stockman & Sharpe LMS cone fundamentals (Stockman & Sharpe, 2000). The gain g parameterises cortical compensation: at $g = 1$ the cortical stage leaves the retinal distortion unchanged; at $g = 2$ it fully cancels the retinal shift; at $g > 2$ cortical processing reverses the direction of the retinal shift, moving the perceived hue past the undistorted colour. The confusion axis is the colour-space direction along which CVD observers cannot distinguish hues from one another. Because $\delta\theta_{\text{Mach}}$ lies entirely along this axis, R+C can only displace colours along this single direction and cannot account for distortions in other colour directions.

Candidate Model 2: two-component cortical (2-comp). We propose a parametric model of the hue distortion pattern measured in § and §??:

$$\delta\theta(\theta; \beta_s, \beta_c) = \beta_s \cos(\theta - 90^\circ) + \beta_c \cos(\theta - \theta_{\text{conf}}) \quad (3)$$

where θ_{conf} is the CVD-subtype confusion axis (16° protan; 150° deutan), derived from Stockman & Sharpe cone fundamentals (Stockman & Sharpe, 2000). β_s parameterises distortion along the S-cone axis ($\theta = 90^\circ$) and β_c parameterises distortion along the individual confusion axis; these two components span the axes of distortion identified by the neural analysis (§, §??). Both signs of β_c are unconstrained; the basis for constraining $\beta_s \geq 0$ is detailed in Supplementary §S1. Model adequacy was assessed by the procedure in §; results are reported in §??.

Inverse fitting

Per-subject distortion parameters (β_s, β_c) were estimated by finding the distortion that, when applied to HC-derived neural representations, best reproduces the observed CVD responses. Three loss families served as the fitting objective.

Behavioral JND ratio (L_γ). For each color pair, the JND ratio $\gamma = \gamma_{\text{CVD}}/\bar{\gamma}_{\text{HC}}$ quantifies how much more difficult discrimination is for a CVD observer relative to the HC mean (§). Two variants enter the fitting: a subtype-targeted pair ratio γ_{subtype} (the color pair predicted to show the largest distortion for each CVD subtype) and an aggregate ratio γ_{all} averaged across all tested pairs. L_γ penalises rank discordance between the model-predicted distortion magnitude at each pair (from Eq. 3) and the observed γ ordering.

Representational $\Delta\text{RDM cosine}$ (L_{RDM}). The between-group geometry difference $\Delta\text{RDM} = \text{RDM}_{\text{CVD}} - \overline{\text{RDM}}_{\text{HC}}$ captures which color pairs drive the CVD–HC representational deviation (§). For a candidate distortion, $\Delta\text{RDM}_{\text{sim}}$ is obtained by applying $\delta\theta$ to the HC mean RDM. The loss is the cosine dissimilarity between the observed and simulated 28-element upper-triangle vectors:

$$L_{\text{RDM}} = \frac{1 - \cos(\Delta\text{RDM}_{\text{sim}}, \Delta\text{RDM}_{\text{obs}})}{2} \quad (4)$$

hV4 LOCO voxel-prediction (L_{LOCO}). The per-hue vulnerability profile $\mathbf{v}_{\text{obs}} \in [0, 1]^8$ records where each CVD participant’s hue interpolation breaks down (§). For a candidate (β_s, β_c) , the simulated profile $\mathbf{v}_{\text{sim}}$ is obtained by re-applying the HC-trained ridge-GCV encoding weights W (§) to the channel design matrix C shifted by the candidate distortion. The loss is the mean squared deviation between the two profiles:

$$L_{\text{LOCO}} = \frac{1}{8} \|\mathbf{v}_{\text{sim}} - \mathbf{v}_{\text{obs}}\|^2 \quad (5)$$

Grid search. Parameters were searched by exhaustive grid over $(\beta_s, \beta_c) \in \langle [-90^\circ, 90^\circ]^2 \rangle$ at 2° resolution ($\langle 8, 281 \rangle$ cells). For R+C (Eq. 2), g was searched over $[0, 3]$. Per-subject loss combinations and parameter selection are described in §.

Parameter selection

Per-subject loss combinations and model class were selected by a three-gate procedure applied to results from a 5-train/2-test HC resample ($N = \langle 300 \rangle$) and a strict 7-fold HC leave-one-out (LOO) using the same atom factories.

Gate 1: Directional precondition. Only loss atoms where the CVD loss value exceeded the HC LOO distribution in the expected direction were admitted. Admission required signed Cohen’s $d \geq +0.5$; atoms failing this gate were excluded from combination.

Gate 2: Boundary saturation. Among admitted combinations, those where $\geq 50\%$ of resample solutions saturated the grid boundary were rejected as indicative of model misspecification.

Gate 3: Held-out test-loss. The primary criterion was the median composite test-loss ($\bar{L}_{\text{test}}$) across the $N = \langle 300 \rangle$ HC resamples. For each resample, the composite loss was evaluated at the training argmin on the two held-out HC, z-normalized relative to the training grid. Lower $\bar{L}_{\text{test}}$ indicates generalization to unseen HC references. The secondary criterion was the test-loss IQR across resamples, measuring resample stability of the generalization estimate. A dedicated neural RDM held-out loss ($\bar{L}_{\text{RDM}}^{\text{LOO}}$; Supplementary §S2) from the strict 7-fold leave-one-out isolates the neural component’s held-out performance from the behavioral weighting in the composite. Supplementary stability indices included the IQR of recovered ($\hat{\beta}_s, \hat{\beta}_c$) (parameter stability), the fraction of resamples recovering the modal 45° bin (mode share), and the strict 7-fold LOO parameter range (Supplementary §S2).

Specificity. The HC false-positive rate was not used as a selection criterion and is reported descriptively in § (Supplementary §S2).

Per-subject selected loss combinations. For sub-08 (deutan), $\gamma_{\text{OY}} + L_{\text{RDM}}^{(V2)}$ satisfied all three gates; for sub-09 (protan), $\gamma_{\text{all}} + L_{\text{RDM}}^{(V1)}$. The LOCO loss family did not enter the winning combination for either participant (Supplementary §S2).

Identifiability and recovery

Identifiability was assessed by three pre-specified checks conducted before parameter reporting: a voxel-level parameter-recovery test, an origin-recovery null test, and HC pseudo-CVD and color-label permutation specificity contrasts (Supplementary §S3). [TODO: revise together with Supplementary S3 once the closure 4-test verdict is finalised.]

Stimulus-space filter

The subject-specific correction filter was derived as the pre-image of the fitted 2-component transform $T(\hat{\beta}_s, \hat{\beta}_c) : \theta \mapsto \theta + \delta\theta(\theta; \hat{\beta}_s, \hat{\beta}_c)$. For each of the 8 displayed hues θ_k , we solved for the input angle $\tilde{\theta}_k$ such that $T(\tilde{\theta}_k) = \theta_k$, yielding a per-hue correction $\delta\theta_k^{\text{filt}} = \tilde{\theta}_k - \theta_k$. A CVD observer viewing $\tilde{\theta}_k$ thus receives, under the model, the same cortical hue representation as a healthy control observer viewing θ_k .

Pre-images were computed by numerical inversion of the forward map (Brent’s-method root-finding); inversion was exact ($|T(\tilde{\theta}_k) - \theta_k| < 0.001^\circ$) for all 8 hues in both CVD participants.

Behavioral tasks

Just-noticeable difference (JND).. Color discrimination thresholds were estimated for eight hue pairs using two-interleaved 1-up/1-down adaptive staircases, each terminating after eight reversals (Levitt, 1971). Pairs spanned three categories: (i) three pairs hypothesized as universally difficult for both CVD types (yellow–purple, blue–purple, red–orange), (ii) two pairs hypothesized as deutan-specific (orange–yellow, yellow–green), (iii) two pairs hypothesized as protan-specific (cyan–magenta, green–blue), and (iv) one control pair (red–cyan). For each pair, the JND ratio $\gamma = \gamma_{\text{CVD}}/\bar{\gamma}_{\text{HC}}$ was computed, where $\bar{\gamma}_{\text{HC}}$ is the HC mean threshold. Ratios $\gamma > 1$ indicate elevated discrimination difficulty relative to HC. These ratios enter the fitting loss as described in §.

Filter evaluation session (CVD only). [TODO: finalise once the Phase 3 design is fixed — specify baseline filter type, run count, and counterbalancing scheme.]

CVD participants completed a second scanning session after the main experiment. Three fMRI runs were acquired under the 2-component-derived stimulus-space filter and three runs under a generic Windows accessibility filter, presented in counterbalanced order. LOCO decoding and LORO classification were repeated on second-session data to assess filter effects on neural color representation. JND behavioral testing was repeated in this session under the same filter conditions, to assess filter effects on discrimination thresholds relative to the unfiltered baseline from Session 1.

[TODO] Planned validation. A two-alternative forced-choice (2AFC) discrimination arm targeting filter-induced perceptual restoration is pre-registered for Phase 3 (data not collected at the time of writing).

Reproducibility

All analyzes were conducted in Python 3.10 (numpy 1.24.3, scipy 1.11.3, scikit-learn 1.3.0, BrainIAK 0.11). Random seeds were fixed (`seed=42`) for all permutation tests and bootstrap resampling. Analysis code and preprocessing scripts are available at https://github.com/Transconnectome/colorBlind_analysis. Anonymised preprocessed data are available on reasonable request to the corresponding author.

Results

Color discrimination is preserved across all ROIs in both CVD cases

A corrective stimulus-space filter is valid only if the visual cortex of a CVD individual retains categorical discrimination among the target hues. We verified this precondition using leave-one-run-out (LORO) cross-validation of the forward encoding model in SRM-aligned space.

Both CVD participants exceeded the 0.125 exact-accuracy chance level at every ROI (Figure 3A). A cross-subject generalization test of the same forward encoding

model (HC-to-HC vs. HC-to-CVD pair-wise cross-decoding, all ROIs pooled) showed no HC–CVD difference (Mann–Whitney U test, 21 HC-to-HC vs. 14 HC-to-CVD pairs, $p = 0.668$, rank-biserial effect size and exact test statistic in Supplementary). The same null held at the within-ROI level at hV4 (Crawford & Howell t -test, both CVD subjects within the HC range, $p = 0.142$). The 8 scanner stimuli are therefore individually distinguishable at the cortical level in both CVD participants, in line with prior work on CVD above-threshold identification (Bosten, 2019; Boehm et al., 2014). The representational substrate on which a corrective filter could act is present.

Color interpolation is impaired in both CVD cases

We next localized where continuous hue interpolation fails, which defines the functional target of the corrective transform, using leave-one-color-out (LOCO) decoding of the forward encoding model. Among the four retinotopic ROIs, only hV4 supported above-chance hue interpolation in HC (adjacent accuracy 0.47 ± 0.05 , mean $\pm$ SEM, $p = 0.044$ under $8! = 40,320$ exact label permutations). V1–V3 did not exceed chance despite preserving cross-run discrimination. This confirms hV4 as the primary interpolation gate, replicating Brouwer and Heeger (2009).

Both CVD participants fell at or below the adjacent-accuracy chance level at hV4 (Figure 3B). Sub-08 reached adjacent accuracy 0.25 (Crawford & Howell $t = -1.58$, $p = 0.082$, n.s., $d_{cc} = -1.71$). Sub-09 reached 0.13 ($t = -2.48$, $p = 0.024$, $d_{cc} = -2.68$). The impairment was not uniform across hues. Both participants showed near-zero adjacent accuracy at hV4 for the S-cone intermediate colors blue, purple, and magenta (adjacent accuracy = 0 for both participants at all three hues; Figure 3C). Per-hue Crawford & Howell tests (one-tailed) reached significance only for blue (both participants $d = 2.20$, $p = 0.042$); purple ($d = 1.02$, $p = 0.19$) and magenta ($d = 1.89$, $p = 0.064$) did not survive the single-case test against the $n = 7$ HC distribution. These per-hue accuracies form each participant’s hue vulnerability profile $\mathbf{v} \in [0, 1]^8$.

Color geometry is distorted at a distinct ROI in each CVD case

To characterize the geometric basis of the LOCO impairment, we compared the pairwise representational structure of CVD and HC participants in a shared SRM-aligned space. For each CVD participant we computed disparity and the per-pair RDM deviation from the HC mean ($\Delta\text{RDM} = \text{RDM}_{\text{CVD}} - \overline{\text{RDM}}_{\text{HC}}$).

The ROI specificity of elevated disparity diverges between the two CVD participants, with V2-dominant disparity in deutan and V1-dominant disparity in protan (Figure 4B). Sub-08 showed significantly elevated disparity specifically at V2 (Crawford & Howell $p = 0.040$), with no significant elevation at V1, V3, or hV4. Sub-09 showed significantly elevated disparity specifically at V1 ($p = 0.007$), with no significant elevation at V2, V3, or hV4. Full test statistics and effect sizes (t , d_{cc}) for these Crawford–Howell comparisons are reported in Supplementary .

The ΔRDM heatmaps (Figure 4A) reveal subject-specific distortion structures. Sub-08 shows elevated distances among S-cone intermediate pairs at V2, and sub-09 shows a broader reorganization of pairwise distances at V1. The ROIs of strongest disparity elevation (V1 for sub-09, V2 for sub-08) differ from the ROI used as the distortion-model loss target (hV4, see the Inverse fitting section). Mean pairwise disparity is an existence statement about absolute geometric perturbation in early visual cortex. hV4 is the cortical stage at which the corrective filter must act, so the ΔRDM enters the loss as a geometric convergence check at that ROI. V1/V2 disparity at upstream ROIs is parallel evidence of distortion rather than a competing fit target.

The retinal-plus-gain model is structurally insufficient

The R+C model is structurally insufficient for both CVD participants. Fitting this model to each CVD participant yielded ceiling saturation in both cases. For sub-08, $\langle 100\% \rangle$ of resample solutions saturated the grid boundary ($g = \langle 3.0 \rangle$). For sub-09, saturation remained high ($\langle 41\% \rangle$, $g = \langle 2.95 \rangle$). Boundary saturation at these rates indicates model misspecification, not a well-constrained optimum ($?$, $?$).

A gain $g > 2$ implies that cortical processing reverses the direction of the retinal

shift past the undistorted hue (Section). Both participants scored below the Ishihara criterion (sub-08: 5/14 plates correct; sub-09: 7/14), confirming residual CVD under standard viewing conditions. A fitted $g > 2$ is therefore internally inconsistent with confirmed CVD status: the model claims overcompensation where the participants demonstrably show impairment. The saturated boundary is a model failure, not a valid parameter estimate.

The structural cause is a degree-of-freedom deficit. The function $\delta\theta = (2 - g) \cdot \delta\theta_{\text{Mach}}$ displaces hues only along the fixed confusion axis. No value of g can account for distortions in directions orthogonal to that axis. R+C is not considered further as a filter candidate.

A common cortical model fits both CVD cases

The 2-component cortical model fit both participants using the same functional form. The deutan and protan cases differ in $(\hat{\beta}_s, \hat{\beta}_c)$, not in model structure. We estimated (β_s, β_c) by casting distortion estimation as an inverse problem (Figure 5): a two-parameter forward model is fit to CVD–HC differences under three loss families, and the fitted transform is analytically inverted to a stimulus-space correction filter.

For sub-08 (deutan, $\theta_{\text{conf}} = 150^\circ$), the selected loss combination $\gamma_{\text{OY}} + L_{\text{RDM}}^{(V2)}$ yielded $(\hat{\beta}_s, \hat{\beta}_c) = (6^\circ, -42^\circ)$. The held-out test-loss was $\bar{L}_{\text{test}} = -2.36$ (IQR = 2.15), lower than the alternative candidate ($\bar{L}_{\text{test}} = -1.14$; Supplementary §S2). The HC-resample parameter IQR was $(8^\circ, 2^\circ)$. The strict 7-fold LOO $\hat{\beta}_c$ range was $[-46^\circ, -38^\circ]$; all values are negative, confirming a stable sign across held-out reference sets.

For sub-09 (protan, $\theta_{\text{conf}} = 16^\circ$), the selected loss combination $\gamma_{\text{all}} + L_{\text{RDM}}^{(V1)}$ yielded $(\hat{\beta}_s, \hat{\beta}_c) = (2^\circ, +24^\circ)$. The held-out test-loss was $\bar{L}_{\text{test}} = -1.54$ (IQR = 1.42). The HC-resample parameter IQR was $(0^\circ, 0^\circ)$, and $\langle 87.7\% \rangle$ of resamples recovered the same 45° bin. The strict 7-fold LOO IQR was $(0^\circ, 0^\circ)$. This position is metric-dependent (Appendix ??).

In both participants, the dominant distortion component is $\hat{\beta}_c$, indicating a

primary distortion along each participant’s confusion axis.

The RDM held-out loss on the strict 7-fold leave-one-out was $\bar{L}_{\text{RDM}}^{\text{LOO}} = 0.594$ for sub-08 and 0.528 for sub-09, beating the no-correction baseline on all 7 held-out folds ((0,0): 1.0; Supplementary §S2). The production fits rank in the top 5–8% of grid combinations, confirming the improvement is non-trivial relative to the no-correction floor. The HC LOO noise ceiling, estimated from split-half HC amplitude reliability (?, ?), was 0.240 (sub-08) and 0.274 (sub-09); this estimate is conservative because CVD amplitude noise is excluded. The production fits recover 52% and 67% of the achievable range toward the noise ceiling, respectively.

The RDM atom in each subject’s winning combination was consistent with the ROI showing significant geometric disparity: V2 for sub-08 ($p = \langle 0.040 \rangle$) and V1 for sub-09 ($p = \langle 0.007 \rangle$, Section). The two criteria were independent: the loss combination was selected by held-out test-loss, and the significant disparity ROI was identified by Crawford–Howell tests on disparity amplitude.

Neural data identifies what behavior cannot

Fitting behavioral (γ) and neural (RDM) loss atoms independently reveals three distinct roles of the neural component.

The neural term captures a protan cortical direction that the behavioral loss cannot resolve. Sub-09’s behavioral-only fit yielded $\hat{\beta}_c \approx +4^\circ$ and did not beat the no-correction baseline ($\Delta L = +0.01$, 4/7 folds). The RDM term captures a signal that the behavioral loss cannot detect for this participant.

The neural term stabilizes the argmin for sub-08. The behavioral-only and combined fits share the same argmin ($6^\circ, -42^\circ$), confirming that the behavioral atom independently locates the deutan direction. Adding the RDM atom reduced the boundary saturation rate from 23% to 9.3%, sharpening the argmin without shifting it.

The neural term also reduces parameter IQR in both participants. For sub-08, adding the neural term reduced the HC-resample IQR from ($18^\circ, 6^\circ$) to ($8^\circ, 2^\circ$) (PCA) and ($10^\circ, 4^\circ$) (SRM). For sub-09, the IQR collapsed from ($6^\circ, 4^\circ$) to ($0^\circ, 0^\circ$) (PCA,

deterministic) and $(0^\circ, 2^\circ)$ (SRM).

The neural-only fit for sub-08 is also non-degenerate ($\hat{\beta}_c = -26^\circ$; Supplementary §S2), independently corroborating the deutan direction identified by the behavioral atom.

Identifiability: mechanism class robust, per-axis value not

Parameter recovery fails across all pre-specified identifiability checks (0/6 significant after FDR correction, BH $\alpha = 0.05$; Supplementary §S3). The voxel-level parameter-recovery test yielded $f_{10^\circ} < 0.30$ for both production candidates. The non-dominant axes ($|\hat{\beta}_s| \leq 6^\circ$ for both participants) fall below the effective parameter uncertainty of $\sim 20^\circ$ – 25° and are not recoverable (Supplementary §S3). The dominant axis ($|\hat{\beta}_c| = 42^\circ$ for sub-08), whose magnitude exceeds this uncertainty, shows partial recovery under GT-consistent synthesis (bias 4.7° ; Supplementary §S3).

Consistent with the partial recoverability of the dominant axis, the sign of $\hat{\beta}_c$ — the distortion direction — is stable across held-out reference sets (Section). Sub-08 consistently returns $\hat{\beta}_c < 0$ across all LOO folds and resamples. Sub-09 consistently returns $\hat{\beta}_c > 0$ under the PCA-basis loss; sign stability under the SRM-basis loss was not verified (Supplementary §S3). The distortion direction is recoverable even when its magnitude is not.

The selected fit lies at the global minimum of each participant’s composite loss surface (Figure 6), but the per-resample argmin distribution makes the identifiability limit explicit: sub-09’s argmin recurs in 87.7% of resamples yet a secondary basin near $(32^\circ, 0^\circ)$ reflects metric-dependent non-identifiability, and sub-08’s argmin disperses along the β_c axis.

Per-subject stimulus-space filter

The per-subject filter $\delta\theta_k^{\text{filt}} = \tilde{\theta}_k - \theta_k$ is the pre-image of the fitted 2-component transform (Section).

For sub-08 (deutan, $\hat{\beta}_s = 6^\circ$, $\hat{\beta}_c = -42^\circ$, $\theta_{\text{conf}} = 150^\circ$), the dominant $\hat{\beta}_c$ term produces a large-amplitude correction that varies asymmetrically across hues (Figure 7).

For sub-09 (protan, $\hat{\beta}_s = 2^\circ$, $\hat{\beta}_c = +24^\circ$, $\theta_{\text{conf}} = 16^\circ$), the smaller $\hat{\beta}_c$ produces a more targeted correction concentrated near the protan confusion axis (Figure 7).

The two correction signatures reflect individual etiologies that a single retinal-family filter cannot capture.

[TODO] Behavioral and neural filter evaluation

Filter efficacy will be evaluated at two levels. (1) Behavioral: same-day 2-alternative forced-choice (2AFC) discrimination of pre-specified hue pairs under the per-subject filter versus a no-filter baseline. (2) Neural: second-session fMRI under the 2-component-derived filter versus a generic Windows accessibility filter (counterbalanced), with LOCO and LORO decoded on second-session data (see the Behavioral tasks section). Results to be added once data analysis is complete.

Discussion

In this study, we tested whether *individual hV4 color geometry* is sufficiently structured to ground a person-specific corrective filter in color vision deficiency. We found that two CVD adults retained 8-class hue discrimination at hV4 but selectively lost continuous-hue interpolation, and that the residual distortion was captured by a 2-component angular-dilation model whose pre-image was exact for all eight hues in both participants (sub-08 mean $|\delta| = 22.2^\circ$; sub-09 mean $|\delta| = 10.2^\circ$). These results suggest that the cortical site at which CVD color information is geometrically distorted is also the site at which it can be inverted into a stimulus-space correction.

The dissociation between LOCO interpolation capacity and SRM ΔRDM geometric distance — 100% vs. 33% concordance with same-day JND thresholds — identifies which cortical quantity is the operative corrective target. Representational geometry (RDM) reflects the metric structure of neural color space; LOCO interpolation reflects whether that structure supports functional prediction of withheld hues. That only LOCO predicts behavioral color confusions is consistent with Brouwer and Heeger (2009), who showed that hV4 — but not early visual areas — supports interpolation to held-out hues in typical observers. The present results extend this

dissociation to these two CVD cases: the same interpolation failure that predicts JND confusion is what the correction filter must address. Geometric distance metrics alone are insufficient to identify it.

The S-cone axis component (β_s) of the fitted distortion field is structurally consistent with prior evidence for cortical color compensation in anomalous trichromats. Emery et al. (2021) reported that CVD perceptual errors concentrate near S-cone confusion loci in behavioral hue-scaling tasks, and Tregillus et al. (2021) showed that V2/V3 BOLD responses in anomalous trichromats include a cortical gain component (g) that partially compensates the retinal cone shift. These individually structured hV4 residuals may reflect cone-class-specific differences in V2/V3 compensation gain that survive into hV4, where the incompletely cancelled L–M signal manifests as a person-specific distortion of the continuous hue code. This suggests that the present 2-component model extends earlier population-mean findings to the individual hV4 level: the residual distortion that survives V2/V3 compensation reaches hV4 as an individually structured pattern. The two CVD participants in fact present qualitatively distinct correction signatures rather than mirror-image realisations of a single mechanism. Sub-08 (deutan) exhibits a wide-amplitude L/M-rotation with large $\hat{\beta}_s = 38^\circ$ and a confusion-axis term $\hat{\beta}_c = -14^\circ$; the R+C diagnostic decomposition recovers a near-zero retinal shift ($\Delta\lambda \approx 2.5$ nm) coupled with a large cortical overshoot ($g = -2.25$), pointing to a cortical-dominant CVD etiology. Sub-09 (protan), by contrast, exhibits a far smaller monotonic rotation along the protan confusion axis (small $\hat{\beta}_s = 6^\circ$, $\hat{\beta}_c = -22^\circ$); the R+C diagnostic recovers a large retinal shift ($\Delta\lambda \approx 19.5$ nm, within the published protan cone-shift range (Tregillus et al., 2021)) and a near-physiological cortical compensation ($g = -1.10$), pointing to a retinal-dominant etiology. This cortical-vs-retinal etiological dissociation is what a population-average retinal model, applied uniformly to all members of a given CVD subtype, fails to capture: in sub-08 the residual distortion is generated downstream of the cone shift and cannot be neutralized by a spectral filter alone, while in sub-09 the etiology is retinal-dominant but the 1-DOF Machado class is structurally non-invertible

at three of this participant’s hues (Appendix , §A.2), so the 2-component class is required for filter realizability in both cases.

All distortion models — 1-DOF Machado, 2-DOF R+C, and 2-component — identify the same two participants as requiring correction, but their prescribed correction vectors diverge in direction (Appendix). This detection–correction dissociation constitutes the study’s primary falsifiable prediction. In a 2AFC experiment comparing the 2-component and retinal-family filter conditions for each participant’s most LOCO-vulnerable hue pairs, only the cortically-grounded 2-component filter should significantly reduce JND. If the retinal-family filter matches or exceeds the 2-component filter in reducing discrimination thresholds, the cortical distortion account would require revision. This preregistered comparison directly tests whether cortical or retinal distortion is the corrective target for CVD color correction.

An alternative reading is that the hV4 LOCO impairment merely reflects reduced V1–V3 input variance rather than cortex-level distortion — that the bottleneck is upstream signal quality and hV4 inherits its deficit passively. The preserved LORO accuracy at hV4 argues against this account: classification and interpolation read out from the same upstream cortical input, so a generic input-variance reduction should degrade both. The observed dissociation instead localises the distortion to the continuous-hue manifold readout, consistent with hV4 as the cortical site of integration over cone-opponent signals (Brouwer & Heeger, 2009; Bannert & Bartels, 2018).

Three considerations bound the proof-of-concept scope of these findings. First, the CVD sample is $N = 2$; single-case statistics (Crawford & Howell, 1998) permit individual-level inference but population-level claims about deutan or protan subtypes require replication with larger samples. Parameter-level bootstrap confidence intervals on $\hat{\beta}_s$ and $\hat{\beta}_c$ were not computed; the inter-subject asymmetry in $\hat{\beta}_s$ (38° vs. 6°) may reflect a robust subtype-level feature, an individual-level idiosyncrasy, or both, and power-adequate replication is required to distinguish these possibilities. The *2-component* label should not be read as confirming a two-mechanism account for both participants. Second, the corrective filter is a stimulus-space prediction derived from a

within-subject neural fit; filter validity is therefore a behavioral question (whether each $\tilde{\theta}_k$ restores the corresponding HC percept), not a within-cohort statistical one. We deliberately do not report fit-landscape permutation p -values for the per-subject $(\hat{\beta}_s, \hat{\beta}_c)$, because such tests address only whether a non-trivial minimum exists on a participant’s vulnerability profile and do not constrain whether the inverted filter restores HC-equivalent perception. Quantitative filter validation is provided by the pre-registered Phase-3 2AFC discrimination arm, which compares pre-image filter, retinal-family filter, and no-filter conditions on each participant’s most LOCO-vulnerable hue pairs. Third, the effective HC pool ($n = 6$ at hV4 after excluding sub-07 for low voxel count) is small enough that the HC $\|\hat{\beta}\|$ leave-one-out distribution we report is a magnitude anchor for the CVD estimates, not a hypothesis test, and external cross-cohort hue-scaling data would be required to arbitrate population-level distortion patterns beyond this proof-of-concept sample.

This pipeline constitutes a proof-of-concept that *individual hV4 color geometry*, measured via fMRI LOCO decoding and inverted through an angular-dilation model, can ground a color correction that is architecturally distinct from population-average retinal approaches. Where Machado-type correction applies a fixed spectral shift to all individuals of a given CVD subtype, the present approach encodes each individual’s hV4 distortion field and inverts it exactly — without assuming population homogeneity. Phase 3 will test whether this cortical-level correction reduces discrimination thresholds where retinal-level correction does not.

References

- Akalin, N. et al. (2025). Improved cvd correction filter using traditional daltonization with mobilenet evaluation. *Symmetry*, *17*, 1046. (Pixel-level Daltonization; CNN used only as automated judge, not neural modeling; non-competitor for neural-filter framing) doi: doi:10.3390/sym17071046
- Andersson, J. L., Jenkinson, M. & Smith, S. (2007). *Non-linear registration aka spatial normalisation* (Technical Report No. TR07JA2). FMRIB.
- Bannert, M. M. & Bartels, A. (2018). Human v4 activity patterns predict behavioral performance in imagery of object color. *Journal of Neuroscience*, *38*(15), 3657–3668. doi: doi:10.1523/JNEUROSCI.2237-17.2018
- Bannert, M. M. & Bartels, A. (2025). Color across human brains: Shared cortical representation of color in healthy observers. *Journal of Neuroscience*, *45*(42).
- Benjamini, Y. & Hochberg, Y. (1995). Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B (Methodological)*, *57*(1), 289–300. doi: doi:10.1111/j.2517-6161.1995.tb02031.x
- Boehm, A. E., MacLeod, D. I. A. & Bosten, J. M. (2014). Compensation for red-green contrast loss in anomalous trichromats. *Journal of Vision*, *14*(13), 19. doi: doi:10.1167/14.13.19
- Bosten, J. M. (2019). The known unknowns of anomalous trichromacy. *Current Opinion in Behavioral Sciences*, *30*, 228–237. doi: doi:10.1016/j.cobeha.2019.10.015
- Brettel, H., Viénot, F. & Mollon, J. D. (1997). Computerized simulation of color appearance for dichromats. *Journal of the Optical Society of America A*, *14*(10), 2647–2655. doi: doi:10.1364/JOSAA.14.002647
- Brouwer, G. J. & Heeger, D. J. (2009). Decoding and reconstructing color from responses in human visual cortex. *Journal of Neuroscience*, *29*(44), 13992–14003. doi: doi:10.1523/JNEUROSCI.3577-09.2009
- Brouwer, G. J. & Heeger, D. J. (2013). Categorical clustering of the neural representation of color. *Journal of Neuroscience*, *33*(39), 15454–15465. doi:

doi:10.1523/JNEUROSCI.2472-13.2013

Chen, P. H., Chen, J., Yeshurun, Y., Hasson, U., Haxby, J. & Ramadge, P. J. (2015).

A reduced-dimension fmri shared response model. In *Advances in neural information processing systems* (Vol. 28, pp. 460–468).

Commission Internationale de L'Eclairage. (1986). *Colorimetry* (2nd ed.). CIE. (Publication No. 15.2)

Conway, B. R., Eskew, R. T., Martin, P. R. & Stockman, A. (2018). A tour of contemporary color vision research. *Vision Research*, 151, 2–6. doi:

doi:10.1016/j.visres.2018.06.009

Crawford, J. R. & Howell, D. C. (1998). Comparing an individual's test score against norms derived from small samples. *The Clinical Neuropsychologist*, 12(4), 482–486. doi: doi:10.1076/clin.12.4.482.7241

Cunningham, J. P. & Yu, B. M. (2014). Dimensionality reduction for large-scale neural recordings. *Nature Neuroscience*, 17(11), 1500–1509. doi: doi:10.1038/nn.3776

Dale, A. M. (1999). Optimal experimental design for event-related fmri. *Human Brain Mapping*, 8(2-3), 109–114. doi:

doi:10.1002/(SICI)1097-0193(1999)8:2/3<109::AID-HBM7>3.0.CO;2-W

Deeb, S. S. (2005). The molecular basis of variation in human color vision. *Clinical Genetics*, 67(5), 369–377. doi: doi:10.1111/j.1399-0004.2004.00343.x

Durnez, J., Moerkerke, B. & Nichols, T. E. (2018). neurodesign: optimal experimental designs for task fmri. *Journal of Open Source Software*, 3(32), 564. doi:

doi:10.21105/joss.00564

Emery, K. J., Volbrecht, V. J., Peterzell, D. H. & Webster, M. A. (2021). Color perception and compensation in color deficiencies assessed with hue scaling.

Vision Research, 183, 1–15. doi: doi:10.1016/j.visres.2021.01.006

Gegenfurtner, K. R. (2003). Cortical mechanisms of colour vision. *Nature Reviews Neuroscience*, 4(7), 563–572. doi: doi:10.1038/nrn1138

Golub, G. H., Heath, M. & Wahba, G. (1979). Generalized cross-validation as a method for choosing a good ridge parameter. *Technometrics*, 21(2), 215–223. doi:

doi:10.1080/00401706.1979.10489751

Gower, J. C. (1975). Generalized procrustes analysis. *Psychometrika*, 40(1), 33–51. doi: doi:10.1007/BF02291478

Hedges, L. V. & Olkin, I. (1985). *Statistical methods for meta-analysis*. Orlando, FL: Academic Press.

Isherwood, Z. J., Joyce, D. S., Parthasarathy, M. K. & Webster, M. A. (2020). Plasticity in perception: insights from color vision deficiencies. *Faculty Reviews*, 9, 8. doi: doi:10.12703/b/9-8

Ishihara, S. (1917). *Tests for colour-blindness*. Tokyo: Kanehara Shuppan.

Jenkinson, M., Bannister, P., Brady, M. & Smith, S. (2002). Improved optimization for the robust and accurate linear registration and motion correction of brain images. *NeuroImage*, 17(2), 825–841. doi: doi:10.1016/s1053-8119(02)91132-8

Khaligh-Razavi, S.-M. & Kriegeskorte, N. (2014). Deep supervised, but not unsupervised, models may explain it cortical representation. *PLOS Computational Biology*, 10(11), e1003915. doi: doi:10.1371/journal.pcbi.1003915

Kriegeskorte, N., Mur, M. & Bandettini, P. (2008). Representational similarity analysis – connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2, 4. doi: doi:10.3389/neuro.06.004.2008

Kuriki, I., Sun, P., Ueno, K., Tanaka, K. & Cheng, K. (2015). Hue selectivity in human visual cortex revealed by functional magnetic resonance imaging. *Cerebral Cortex*, 25(12), 4869–4884. doi: doi:10.1093/cercor/bhv194

Levitas, D., Hayashi, S., Vinci-Booher, S., Heinsfeld, A., Bhatia, D., Lee, N., . . . Pestilli, F. (2024). ez bids: Guided standardization of neuroimaging data interoperable with major data archives and platforms. *Scientific Data*, 11, 179. doi: doi:10.1038/s41597-024-02959-0

Levitt, H. (1971). Transformed up-down methods in psychoacoustics. *The Journal of the Acoustical Society of America*, 49(2), 467–477. doi: doi:10.1121/1.1912375

Machado, G. M., Oliveira, M. M. & Fernandes, L. A. F. (2009). A physiologically-based model for simulation of color vision deficiency. *IEEE Transactions on*

- Visualization and Computer Graphics*, 15(6), 1291–1298. doi: doi:10.1109/TVCG.2009.113
- Maes, F., Collignon, A., Vandermeulen, D., Marchal, G. & Suetens, P. (1997). Multimodality image registration by maximization of mutual information. *IEEE Transactions on Medical Imaging*, 16(2), 187–198. doi: doi:10.1109/42.563664
- Neitz, J. & Neitz, M. (2011). The genetics of normal and defective color vision. *Vision Research*, 51(7), 633–651. doi: doi:10.1016/j.visres.2010.12.002
- Nichols, T. E. & Holmes, A. P. (2002). Nonparametric permutation tests for functional neuroimaging: a primer with examples. *Human Brain Mapping*, 15(1), 1–25. doi: doi:10.1002/hbm.1058
- Parkes, L. M., Marsman, J. B. C., Oxley, D. C., Goulermas, J. Y. & Wuerger, S. M. (2009). Multivoxel fmri analysis of color tuning in human primary visual cortex. *Journal of Vision*, 9(1), 1.1–1.13. doi: doi:10.1167/9.1.1
- Patterson, E. J., Mastey, R. R., Kuchenbecker, J. A., Rowlan, J., Neitz, J., Neitz, M. & Carroll, J. (2022). Effects of color-enhancing glasses on color vision in congenital red–green color deficiencies. *Optics Express*, 30(17), 31182. doi: doi:10.1364/OE.461425
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., . . . Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- Peirce, J., Gray, J. R., Simpson, S., MacAskill, M., Höchenberger, R., Sogo, H., . . . Lindeløv, J. K. (2019). Psychopy2: Experiments in behavior made easy. *Behavior Research Methods*, 51(1), 195–203. doi: doi:10.3758/s13428-018-01193-y
- Schütt, H. H., Kipnis, A. D., Diedrichsen, J. & Kriegeskorte, N. (2023). Statistical inference on representational geometries. *eLife*, 12, e82566. doi: doi:10.7554/eLife.82566
- Shapley, R. & Hawken, M. J. (2011). Color in the cortex: single- and double-opponent cells. *Vision Research*, 51(7), 701–717. doi: doi:10.1016/j.visres.2011.02.012
- Somers, L. P., Franklin, A. & Bosten, J. M. (2024). Empirical tests of the effectiveness

- of EnChroma multi-notch filters for enhancing color vision in deuteranomaly. *Vision Research*, 218, 108390. doi: doi:10.1016/j.visres.2024.108390
- Stockman, A. & Sharpe, L. T. (2000). The spectral sensitivities of the middle- and long-wavelength-sensitive cones derived from measurements in observers of known genotype. *Vision Research*, 40(13), 1711–1737. doi: doi:10.1016/S0042-6989(00)00021-3
- Tregillus, K. E. M., Isherwood, Z. J., Vanston, J. E., Engel, S. A., MacLeod, D. I. A., Kuriki, I. & Webster, M. A. (2021). Color compensation in anomalous trichromats assessed with fmri. *Current Biology*. doi: doi:10.1016/j.cub.2021.01.039
- Wang, L., Mruczek, R. E., Arcaro, M. J. & Kastner, S. (2015). Probabilistic maps of visual topography in human cortex. *Cerebral Cortex*, 25(10), 3911–3931. doi: doi:10.1093/cercor/bhu277
- Webster, M. A. (2015). Visual adaptation. *Annual Review of Vision Science*, 1, 547–567. doi: doi:10.1146/annurev-vision-082114-035509
- Wells, W. M., Viola, P. A., Atsumi, H., Nakajima, S. & Kikinis, R. (1996). Multi-modal volume registration by maximization of mutual information. *Medical Image Analysis*, 1(1), 35–51. doi: doi:10.1016/S1361-8415(01)80004-9

Appendix A: Alternative Model Comparison

The main text adopts a 2-component cortical model (Section , fit under the V4 LOCO neural-primary loss L_{fit} of Equation ??). Here we report two cone-shift baselines as reference comparisons:

- **Machado (1-DOF)** — a single retinal cone-spectral shift $\Delta\lambda$ Machado et al., 2009.
- **Retinal + Cortical (R+C, 2-DOF)** — retinal shift $\Delta\lambda$ plus a cortical compensation gain g Tregillus et al., 2021.

A.1 Fit results (LOCO Spearman ρ argmax, historical reference)

Under the LOCO- ρ argmax criterion used in earlier drafts, the Machado model fit Sub-09 at $\Delta\lambda = 13.5$ nm with $\rho_{V4} = 0.76$ and Sub-08 at $\rho_{V4} = 0.62$ (lower fit quality); the R+C model did not improve on either single-component model in either subject. These fits are reported as historical reference for the main-text 2-component fit (Section); the present L_{fit} loss (Equation ??) and LOCO- ρ argmax need not agree, since L_{fit} optimizes a combination of per-color V4 vulnerability MSE, Spearman rank ordering, ΔRDM cosine, and a smoothness regularizer rather than the rank-order monotonicity of ρ at hV4 alone.

R+C decomposition as a diagnostic. Although the R+C 2-stage model is not adopted as the filter form (see Per-subject stimulus-space filter), the R+C parameter fit at hV4 under L_{fit} is retained as a per-subject etiological diagnostic. For sub-08, the R+C decomposition recovers $\Delta\lambda = 2.5$ nm (near-zero retinal shift) with $g = -2.25$ (large cortical overshoot), consistent with a cortical-dominant CVD. For sub-09, the R+C decomposition recovers $\Delta\lambda = 19.5$ nm (within the published protan cone-shift range) with $g = -1.10$ (near-physiological cortical compensation, no overshoot), consistent with a retinal-dominant CVD. This per-subject cortical-vs-retinal asymmetry is reported as a mechanistic interpretation of the 2-component fit, not as an independent test.

A.2 Inverse-mapping (pre-image) behavior

The angular form of the 2-component model is bijective by construction: a stimulus-space pre-image $\theta + \delta\theta$ exists for any target hue. The 1-DOF Machado fit for sub-09, by contrast, collapses three hues (green 135° , cyan 180° , blue 225°) onto a single pre-image angle ($\sim 127^\circ$), which mathematically prevents exact inversion at those positions. This arc-compression failure is a structural property of the 1-DOF retinal-shift model and is independent of the fitting criterion; it disqualifies the Machado baseline for the stimulus-space filter task of Section .

A.3 R+C model

The R+C model is retained here as a 2-DOF baseline. Because the cortical gain term g acts as a uniform amplitude scaling without generating the angular distortion captured by the 2-component model's $\beta_c \cos(h - \theta_{\text{conf}})$ term, R+C does not match the observed hV4 LOCO vulnerability pattern beyond what the retinal $\Delta\lambda$ term alone provides; this conclusion is consistent across loss criteria.

Supplementary Methods

S1. Confound Regression and Temporal Filtering

Within-run head motion was estimated via rigid-body realignment. No temporal filtering or confound regression was applied; slow drift was modeled via linear per-run regressors in the general linear model.

S2. Quality Control

Registration quality was evaluated by computing the intersection between the Wang Atlas ROIs (V1, V2, V3, hV4) and the BOLD brain mask in MNI space. Across 10 subjects, mean ROI coverage was 84.3% (SD = 21.7%), and GLM valid ratio (percentage of ROI voxels with reliable stimulus-evoked responses) was 99.6%. We evaluated all subjects' data suitable for downstream analysis, though sub-07 showed reduced coverage (30.8%) due to individual anatomical variability.

S3. Voxel Counts Range

Voxel counts after ROI definition and quality filtering are reported in the main Methods section. Detailed per-subject counts are available in the online repository.

S4. Consistency Test: PCA and Procrustes (TBA)

Additional validation analyzes comparing SRM with PCA-based dimensionality reduction and norm-constrained Procrustes alignment will be reported following completion of the full dataset.

S5. Mean Activation Analysis

We conducted mean voxel response amplitudes analysis to identify whether SRM-based differences rely on the activation. We checked mean amplitudes across all colors and for each color.

S6. Statistical Analysis

All tests used a significance threshold of $\alpha = 0.05$ with False Discovery Rate correction (Benjamini & Hochberg, 1995); $q = 0.05$ applied to RDM color-pair comparisons (28 pairs per participant). Individual CVD participants were compared to the HC distribution using Crawford and Howell (1998) modified t-tests:

$$t^* = \frac{x_{\text{CVD}} - \bar{x}_{\text{HC}}}{\text{SD}_{\text{HC}} \times \sqrt{(n+1)/n}}, \quad df = n - 1 = 6 \quad (6)$$

with one-tailed p-values testing the directional hypothesis that CVD disparity exceeds HC. Effect sizes were quantified via Hedges' g . Group-level permutation tests (10,000 iterations) were used for HC-CVD comparisons by permuting subject labels; color-level permutation tests (1,000 iterations) were used for LOCO/LORO validation by permuting color labels within subjects.

S7. Dimensionality Selection (K-Selection)

The reduced dimension k for SRM was determined via leave-one-subject-out (LOSO) cross-validation based on within-subject representational dissimilarity matrix (RDM) reliability (Cunningham & Yu, 2014). For each candidate k (range: 1–8) and fold, reliability was quantified as the Spearman correlation between RDMs computed from even runs (2,4,6) versus odd runs (1,3,5) of the held-out participant in SRM space. We selected k maximizing mean reliability across seven folds. The selected values and their associated cross-subject RDM Spearman correlations were:

- V1: $k = 4$ (cross-subject RDM correlation = 0.60)
- V2: $k = 4$ (cross-subject RDM correlation = 0.57)
- V3: $k = 3$ (cross-subject RDM correlation = 0.55)
- hV4: $k = 3$ (cross-subject RDM correlation = 0.32)

S8. Generalized Cross-Validation for Ridge Regression

The regularization parameter α for ridge regression was selected via Generalized Cross-Validation (GCV; Golub et al., 1979), implemented as `sklearn.linear_model.RidgeCV` with `gcv_mode='auto'` (Pedregosa et al., 2011). GCV provides a computationally efficient approximation to leave-one-out cross-validation for selecting α . The optimal α minimizes the GCV score:

$$\text{GCV}(\alpha) = \frac{1}{n} \|y - \hat{y}\|_2^2 / \left(1 - \frac{\text{tr}(H)}{n}\right)^2 \quad (7)$$

where X denotes the design matrix, y is the observed response vector, $\hat{y}$ is the ridge-predicted response, $H = X(X^T X + \alpha I)^{-1} X^T$ is the hat matrix, and n is the number of training samples.

S9. Leave-One-Out Consistent Disparity Estimation

To quantify HC-CVD disparity with unbiased references, we used a leave-one-out (LOO) consistent procedure: for each fold i (leaving out HC participant i), a reference pattern R_{-i} was computed as the mean of the remaining six HC participants' aligned data. Both the held-out HC_i and both CVD participants were compared against this same reference R_{-i} , yielding disparity scores $d(\text{HC}_i, R_{-i})$ and $d(\text{CVD}_j, R_{-i})$. This ensures HC and CVD disparities are measured against identical references within each fold. Individual CVD disparity significance was assessed via Crawford and Howell (1998) modified t-tests comparing each CVD participant's mean disparity (averaged across 7 folds) to the HC distribution ($n = 7$, $df = 6$).

S10. Cross-Validation Procedures

Leave-One-Run-Out (LORO)

LORO evaluated whether individual colors could be distinguished from voxel response amplitudes, testing the consistency of neural color encoding across runs. For each of six runs, the corresponding response matrices (B , C) were held out as the test

set, W was trained on the remaining five runs. Both color decoding and voxel response prediction were evaluated on the held-out run.

Leave-One-Color-Out (LOCO)

LOCO tested whether geometric relationships of circular hue space were preserved in neural representation by evaluating interpolation to held-out colors. For each color, all six run response matrices of the target color were held out. W was trained on the remaining seven colors, and the model predicted either the held-out color from voxel responses or voxel responses from the held-out color’s channel outputs.

Leave-One-Subject-Out (LOSO)

LOSO evaluated zero-shot generalization to held-out participants. As LOSO requires cross-subject prediction, SRM-transformed data were used to align voxel spaces across participants. For each HC participant, W was trained on the remaining HC participants’ shared-space data and tested on the held-out participant. This assessed whether the HC group-derived common space supported generalization to novel individuals.

S11. Evaluation Metrics

Color Decoding Metrics

Color decoding was evaluated using (1) mean absolute error (MAE) between predicted and true hue angles (chance = 90° , expected from a uniform random predictor on a circle), and (2) categorization accuracy with $\pm 22.5^\circ$ bins around each of the eight colors (chance = $12.5\% = 1/8$).

Voxel Prediction Metrics

Voxel response prediction was evaluated using Pearson correlation between predicted and observed response vectors, computed across all voxels and colors in the test set.

Statistical Significance

Statistical significance of group comparisons (HC vs CVD) and individual CVD analyzes was assessed using nonparametric permutation tests (10,000 iterations; Nichols & Holmes, 2002).

S12. Shared Response Model Residuals

Each participant’s residuals in the SRM model reflect the extent to which their data are not captured by the common structure, including inherent geometric differences. The orthogonal constraint ($W_i^T W_i = I_k$) helps preserve geometric relationships between colors—distances and angles—in the common space.

Additionally, global between-subject RDM similarity was quantified using Spearman correlation across all 28 color pairs.

S13. Image Orientation Initialization

Image orientation was first initialized according to the scanner-defined obliquity information in the NIfTI header. The mutual-information cost function was then optimized to maximize the statistical dependence between T1w and BOLD intensity distribution within the overlapping region.

S14. Procrustes Alignment Benefits

This preprocessing reduces measurement noise across runs while preserving color representational geometry, thereby improving the quality of subsequent SRM fitting.

S16. Comparison with Retinal-Family Distortion Models

Two retinal-family models were evaluated alongside the 2-component model by the same grid-search and permutation procedure (see Section).

Models. The 1-DOF Machado model Machado et al., 2009 shifts the L/M cone spectral sensitivity by $\Delta\lambda$ nm, yielding a hue-confusion pattern characteristic of the CVD subtype. The 2-DOF retinal+cortical R+C model Tregillus et al., 2021 augments

Machado with a cortical gain g : $r'_g = r_{g,\text{base}} + (1 + g)(r_{g,\text{ret}} - r_{g,\text{base}})$, where $g = 0$ recovers the retinal-only model and $g = -1$ is exact cortical compensation.

Fits (Table 1). Under the historical LOCO- ρ argmax criterion (40,320 label permutations), reported here for reference only, retinal-family fits at hV4 yield Sub-08 R+C ($\Delta\lambda = 2.0$ nm, $g = +2.25$, $\rho = 0.857$) and Sub-09 1-DOF Machado ($\Delta\lambda = 13.5$ nm, $\rho = 0.762$); the latter shift magnitude falls within the established range for moderate–severe protanomaly (Neitz & Neitz, 2011). The large positive cortical gain for Sub-08 ($g = +2.25$, indicating >100% overshoot beyond the exact Tregillus compensation level of $g = -1$; Tregillus et al., 2021) is not physiologically interpretable as a gain parameter and is treated as an empirical characterization of residual hV4 structure. The cross-subject model-class divergence (R+C for Sub-08; 1-DOF Machado for Sub-09) prevents a common-axis comparison of the two cases’ distortion mechanisms.

Correction-vector divergence. Despite comparable LOCO fit quality ($\rho = 0.857$ for R+C vs. 0.881 for 2-component), the R+C and 2-component correction vectors for Sub-08 have cosine similarity -0.18 and sign agreement only 3/8, confirming that detection converges while correction directions diverge at the cortical level.

Δ RDM cross-model convergence. Using R+C parameters for Sub-09 ($\Delta\lambda = 19.5$ nm, $g = -1.10$), an independent Δ RDM test reaches $p = 0.026$ (40,320 cosine-similarity permutations on crossnobis distances), providing cross-model corroboration for the Sub-09 distortion signal. Retinal-family optimization directly against Δ RDM fails for both CVD subjects; LOCO is the more discriminating phenotype and Δ RDM serves as a convergence cross-check only.

Arc-compression and LOCO vulnerability. At Sub-09’s severity, Machado arc-compression ($\Delta\lambda = 13.5$ nm) collapses three hues (c4 green, c5 cyan, c6 blue) to a single perceived angle in opponent hue space. These three collapsed hues fall within Sub-09’s LOCO vulnerability cluster (c5/c6 cyan/blue), confirming that arc-compression eliminates correction capacity precisely where Sub-09 requires it most.

S17. HC leave-one-out magnitude-anchor distribution

Each of $n = 6$ effective HC participants was refit under the canonical V4 LOCO loss L_{fit} (Equation ??) with the remaining HC pool as reference, yielding an HC leave-one-out distribution of fitted-parameter magnitudes $\|\hat{\beta}\|$ ranging 26.3° to 49.2° (mean 40.1°). Sub-07 was excluded because its hV4 ROI contained only 16 voxels (versus 67–70 for the remaining HC), which precluded a stable forward-encoder LOCO target vector at this ROI; the exclusion arises in filter-fit space and is unrelated to the SRM dimensionality choice. This range is reported as a magnitude anchor for the CVD estimates (sub-08 $\|\hat{\beta}\| = 40.5^\circ$, falling near the HC mean; sub-09 $\|\hat{\beta}\| = 22.8^\circ$, below the HC minimum), not as a hypothesis test of CVD-versus-HC specificity.

S18. Effect Sizes for Single-Case Comparisons

Crawford–Howell single-case comparisons used the modified t Crawford & Howell, 1998; the corresponding case-control effect size is $d_{cc} = (x_{\text{case}} - \bar{x}_{\text{HC}})/s_{\text{HC}} = t \cdot \sqrt{(n+1)/n}$ where n is the HC sample size used in the test. Mann–Whitney U tests are reported with rank-biserial correlation r_{rb} as effect size. Full per-test statistics ($t, U, p, d_{cc}, r_{\text{rb}}$) for every Crawford–Howell and Mann–Whitney comparison reported in the main text are tabulated in the online repository (`statistics/effect_sizes.csv`).

Table 1

All distortion model fits at hV4 LOCO for both CVD participants. Retinal-family fits are reported under the historical LOCO- ρ argmax criterion (Section A.1 of Appendix A); ρ is Spearman correlation between simulated and observed per-color vulnerability.

Subject	Model	DOF	Param. 1	Param. 2	ρ
Sub-08 (deutan)	2-component	2	$\beta_s = 38^\circ$	$\beta_c = -14^\circ$	0.881
Sub-09 (protan)	2-component	2	$\beta_s = 6^\circ$	$\beta_c = -22^\circ$	0.690
Sub-08 (deutan)	R+C	2	$\Delta\lambda = 2.0$ nm	$g = +2.25$	0.857
Sub-09 (protan)	Machado	1	$\Delta\lambda = 13.5$ nm	—	0.762




Figure 1. Experimental paradigm and analysis pipeline. (A) Stimulus set: 8 isoluminant CIE $L^*a^*b^*$ hues at 45° angular spacing ($L^* = 75$, chroma = 40; red–magenta). (B) Retinotopic ROIs: V1, V2, V3, hV4 (standard retinotopic mapping). (C) Analysis pipeline. Stage A: two-stage GLM single-trial β amplitudes, Procrustes-aligned across runs, projected into an SRM space trained on HC participants only. Decoding: LORO (discrimination), LOCO (interpolation). CVD characterisation: 2-component cortical model with S-cone-axis (β_s) and confusion-axis (β_c) components. Pre-image yields per-subject stimulus-space filter. HC: $n = 7$ (sub-01–07); deutan: sub-08; protan: sub-09.




Figure 2. Forward encoding model schematic. Stage 1: each displayed hue is encoded as channel activations by six half-wave rectified squared-cosine basis functions, with preferred hues spaced at 60 degrees (Brouwer & Heeger 2009). Stage 2 (training): a per-voxel weight matrix W is fit by ridge regression with GCV-selected regularization, mapping channels to voxel responses. Stage 3 (testing): for a held-out voxel pattern, channel activations are reconstructed as $\hat{c} = Wx$. Stage 4: the decoded hue is recovered as the angle whose basis vector has the highest Pearson correlation with the reconstructed channel pattern. The same encoder is used for cross-run color-discrimination decoding (Section) and cross-color interpolation decoding (Section); the test set differs but the basis, weights, and correlation-based selection are identical.




Figure 3. Color discrimination is preserved but interpolation is selectively impaired at hV4 in CVD. (A) LORO discrimination accuracy (forward encoding model, 8-class, SRM-aligned), V1–hV4. HC mean $\pm$ SEM ($n = 7$, gray); sub-08 (orange square), sub-09 (teal triangle). Chance = $1/8$. *(B)* LOCO adjacent accuracy (± 1 hue step; chance = $3/8$). *: sub-09 below HC at hV4 (Crawford & Howell $t = -2.48$, $p = 0.024$; exploratory single-case finding); sub-08 $t = -1.58$, $p = 0.082$, n.s. *(C)* Per-hue adjacent accuracy at hV4, HC bars $\pm$ SEM. Per-hue Crawford & Howell tests (one-tailed) were significant only for blue (both participants $d = 2.20$, $p = 0.042$); purple ($d = 1.02$, $p = 0.19$) and magenta ($d = 1.89$, $p = 0.064$) were n.s. Joint pattern in (A) and (B–C) defines the filter precondition and target, respectively.





Figure 4. Each CVD subject shows significantly elevated color representation disparity at a distinct ROI. (A) $\Delta RDM = RDM_{CVD} - \overline{RDM}_{HC}$ in SRM-aligned space at each subject's primary ROI (sub-08 V2, sub-09 V1). Warm: CVD distances larger than HC mean; cool: smaller. **(B)** Mean pairwise correlation distance (disparity) per subject per ROI. HC band: group mean $\pm$ 1 SD ($n = 7$); individual HC LOO as dots. Crawford & Howell tests (sub-08 V2: $p = 0.040^*$; sub-09 V1: $p = 0.007^{**}$, exploratory single-case). (A) shows the directional pattern of pairwise deviations. (B) tests the absolute amplitude of pairwise distances against the HC LOO distribution. Idiosyncratic ROI specificity (V2 deutan; V1 protan) is inconsistent with a shared group-level gain mechanism.



Schematic; glyphs illustrative

Figure 5. Inverse-problem workflow for personalised stimulus-space filter

design. (1) A two-parameter forward model expresses the per-hue distortion as $\delta\theta(\theta) = \beta_s \cos(\theta - 90^\circ) + \beta_c \cos(\theta - \theta_{\text{conf}})$: an S-cone-axis component (β_s) plus a family-specific confusion-axis rotation (β_c ; $\theta_{\text{conf}} = 16^\circ$ protan, 150° deutan). (2) Three loss families quantify CVD–HC differences — behavioural JND ratios (γ), neural representational ΔRDM cosine, and hV4 forward-encoding LOCO; only the LOCO loss couples to the forward encoder. (3) A grid search over (β_s, β_c) minimises the z -scored composite to the selected fit $\hat{\beta}$. (4) The fitted forward map is analytically inverted: for each displayed hue θ , the corrected stimulus θ_{pre} solves $\theta_{\text{pre}} + \delta\theta(\theta_{\text{pre}}) = \theta$. (5) The per-subject filter is $\delta\theta_{\text{filter}}(\theta) = \theta_{\text{pre}} - \theta$. Glyphs are illustrative; the fitted loss landscape is shown in Figure 6.




Figure 6. Per-subject loss landscape and parameter uncertainty. z -scored composite loss over the (β_s, β_c) grid for each CVD participant’s selected loss combination (sub-08: $\gamma_{\text{OY}} + L_{\text{RDM}}^{(V2)}$; sub-09: $\gamma_{\text{all}} + L_{\text{RDM}}^{(V1)}$), reconstructed on the full 7-HC pool. Lower (yellow) indicates better fit. Red star: selected fit (argmin) — sub-08 $(\hat{\beta}_s, \hat{\beta}_c) = (6^\circ, -42^\circ)$; sub-09 $(2^\circ, +24^\circ)$. White points: per-resample argmins over $N = 300$ HC 5-train/2-test resamples, showing parameter uncertainty. The selected fits are descriptive low-dimensional embeddings of each participant’s distortion, not physiological point estimates (this section).




Figure 7. Per-subject stimulus-space correction filter. For each CVD participant, the displayed stimulus (*Original*) and the corrected stimulus (*Filtered*) are shown for the eight scanner hues. Left block: sub-08 (deutan), $(\hat{\beta}_s, \hat{\beta}_c) = (6^\circ, -42^\circ)$. Right block: sub-09 (protan), $(\hat{\beta}_s, \hat{\beta}_c) = (2^\circ, +24^\circ)$. The filtered stimulus is the exact numerical pre-image of the fitted 2-component transform; the per-hue correction $\delta\theta_{\text{filter}} = \theta_{\text{pre}} - \theta$ is annotated beneath each *Filtered* swatch (8/8 pre-images exact, residual $< 0.01^\circ$ for both subjects). Rendering coordinate space: STIM_LAB CIELab (project convention). Quantitative behavioural validation of filter efficacy is deferred to the Phase-3 2AFC arm and will be added once data are collected.